

Supervised Learning


Vishal Patel

Spring 2022



THE NEW YORK TIMES BESTSELLER

THINKING,
FAST AND SLOW



DANIEL
KAHNEMAN

WINNER OF THE NOBEL PRIZE IN ECONOMICS

"[A] masterpiece. . . This is one of the greatest and most engaging collections of insights into the human mind I have read." —WILLIAM EASTERLY, *Financial Times*

Course Outline

1. Introduction

2. The Data Science Process

3. Supervised Learning

4. Unsupervised Learning

5. The Grunt Work

6. Wrap Up

Supervised Learning

1 Train (a model)

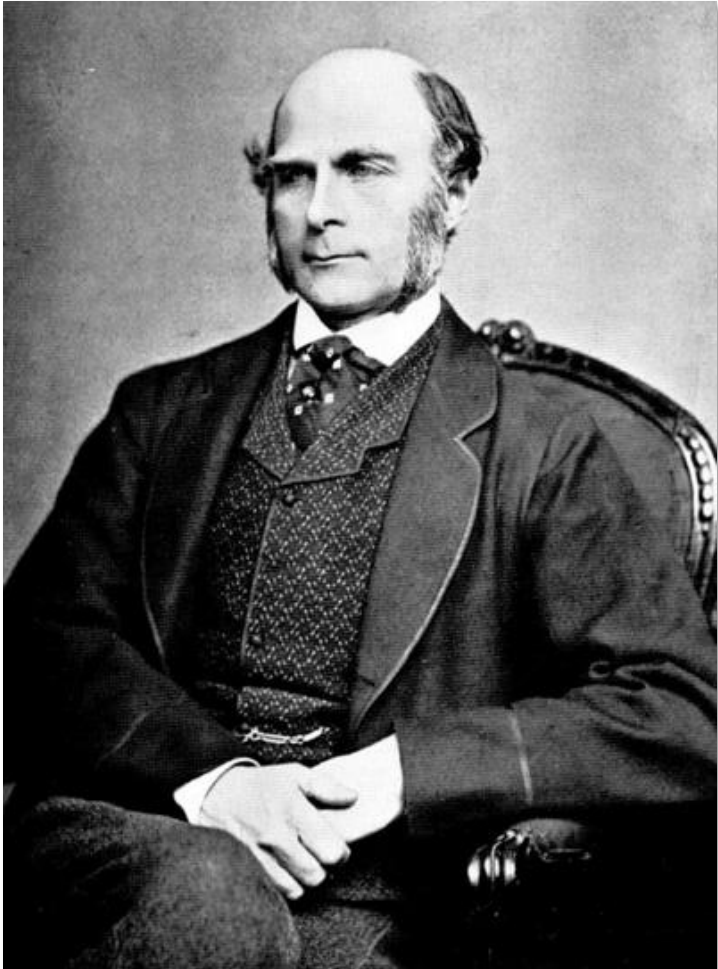
2 Make Predictions

Observations + Labels

Ground Truth
↑

Observations → Predictions

Regression to Mediocrity



Sir Francis Galton

1822 – 1911

An English Victorian era statistician, progressive, polymath, sociologist, psychologist, anthropologist, eugenicist, tropical explorer, geographer, inventor, meteorologist, proto-geneticist, and psychometrician.


Regression to Mediocrity


A larger-than-average parent tends to produce
a larger-than-average child,
but the child is **likely** to be *less* large than the parent
in terms of its relative position within its *own* generation.

Regression to the *mean*

RATE OF REGRESSION IN HEREDITARY STATURE.

Fig. (a)





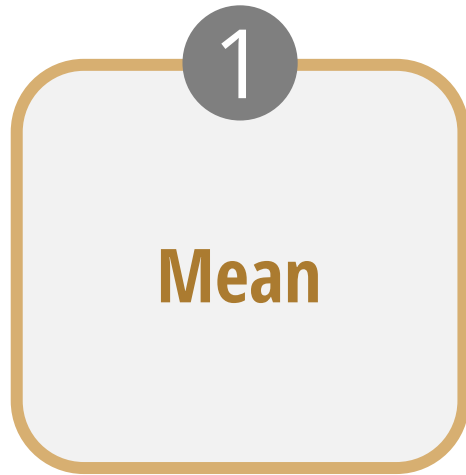
- Here we selected the group of students whose grades were worse than average in the first test.
- The fact that they did poorly in the first test means that both skill and luck (random chance) were NOT in their favor.
- Now during the second test, we may expect them to be equally not skillful, but we should not expect all of them to be equally unlucky.
- Hence, we should predict that their score on the second test would be closer to the mean than before.
- Similarly, in Kahneman's example, if a cadet did something exceptional, his/her next attempt is unlikely to be as good (whether he/she was praised or not.)
- Your children can be expected to be less exceptional (for better or worse) than you are. A baseball player's batting average in the second half of the season can be expected to be closer to the mean (for all players) than his batting average in the first half of the season. And so on.
- The key word here is "expected".

Linear Regression

Linear Regression

- **Variable:** A quantity that may vary across observations (either measurements taken across different times or across different subjects, e.g., people).
- When we fit a linear model, we assume (or hope) that one variable (e.g., y) do not vary randomly, but varies as a straight-line function of another variable (e.g., x). In other words, y is **dependent** on x .
- How do we measure this dependence?
- **Variance:** A measure of the amount of variability in a variable, and it's defined as average squared deviations (fluctuations) from its mean.
 - Alternatively, we can measure variability in terms of standard deviation, which is defined as the square root of variance.
- The goal of a liner model is to find out **how much** of the variation in y (fluctuations from its mean) can be explained by variation in x (fluctuations from its mean).

A linear regression model can be estimated based on only **three** statistics:



$$r_{xy}$$

If we know the **correlation coefficient**, we know the extent to which **fluctuations** of one variable (x) from its **mean** can be used to predict the **fluctuations** of other variable (y) from its **mean**.

Correlation Coefficient

$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

It measures the strength of the linear relationship
between y and x on a scale of -1 to $+1$.

In order to calculate the correlation coefficient,
we should first **standardize** the variables
by taking out the mean and dividing by standard deviation.

How much does each point
fluctuate from its mean

Z Score

x_i^* =

$$\frac{x_i - AVERAGE(x)}{STDEV(x)}$$

... compared to
how much this variable fluctuates overall.

$$x_i^* = \frac{x_i - AVERAGE(x)}{STDEV(x)}$$



$$x_i^* = \frac{(x_i - \bar{x})}{\frac{1}{n} \sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}}$$

$$y_i^* = \frac{y_i - AVERAGE(y)}{STDEV(y)}$$



$$y_i^* = \frac{(y_i - \bar{y})}{\frac{1}{n} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

$$r_{xy} = \frac{1}{n} \sum_{i=1}^n x_i^* y_i^*$$

$$r_{xy} = \frac{1}{n} \sum_{i=1}^n x_i^* y_i^*$$



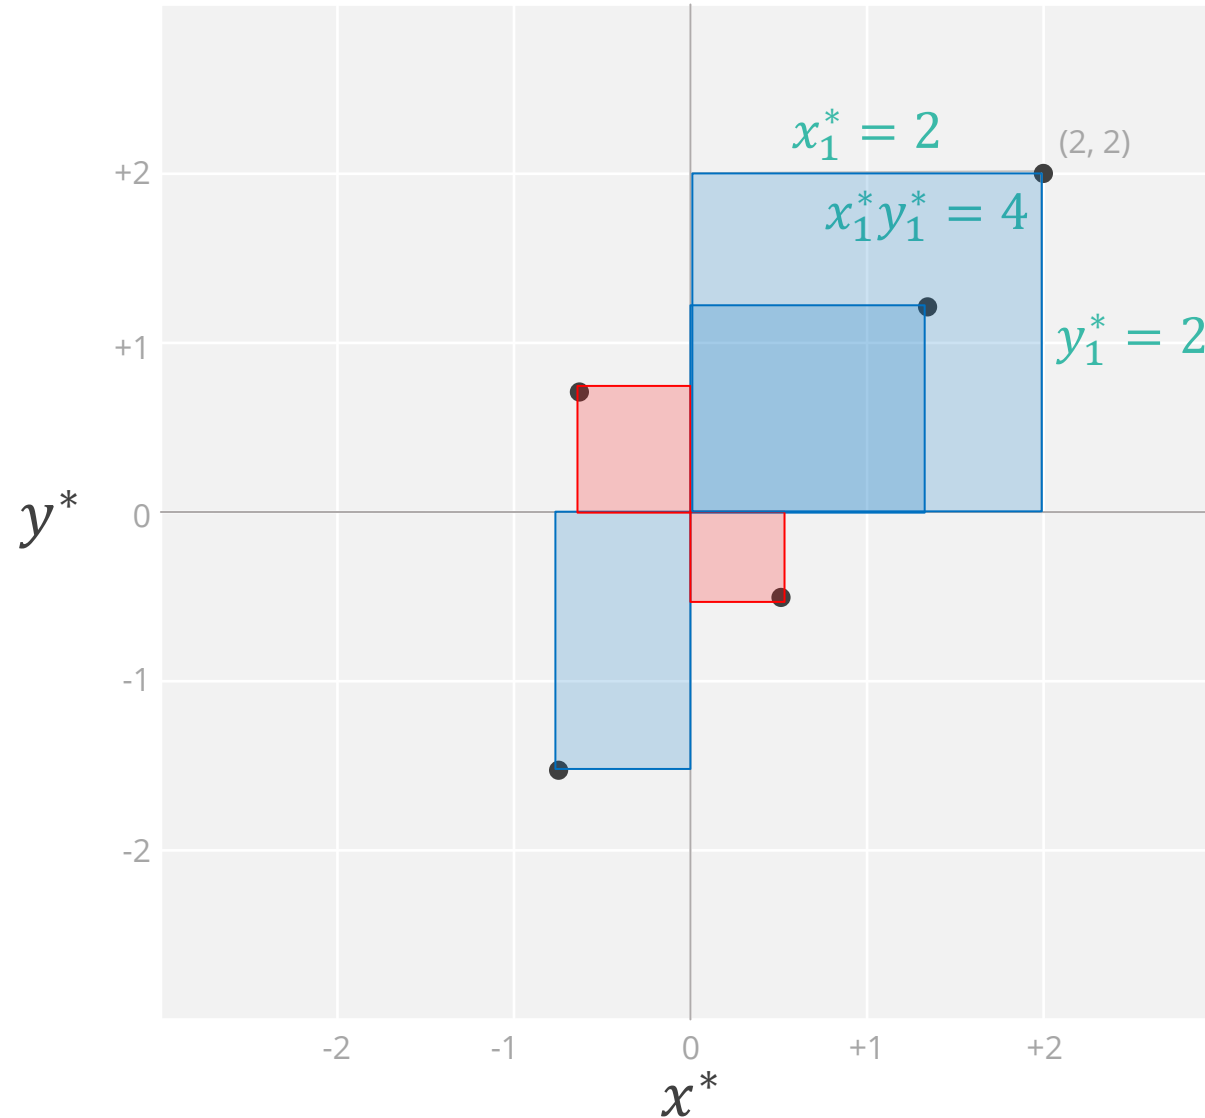
$$r_{xy} = \frac{(x_1^* y_1^* + x_2^* y_2^* + \cdots + x_n^* y_n^*)}{n}$$

The **correlation coefficient** is equal to the **average product of the standardized values** of the two variables.

$$x_i^* = \frac{(x_i - \bar{x})}{\frac{1}{n} \sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}}$$

$$y_i^* = \frac{(y_i - \bar{y})}{\frac{1}{n} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

$$r_{xy} = \frac{(x_1^*y_1^* + x_2^*y_2^* + x_3^*y_3^* + x_5^*y_5^* + x_5^*y_5^*)}{5}$$



The sum of those five rectangles would add up to a positive number, hence we would get a positive correlation coefficient between x and y .

The correlation coefficient
measures the strength of the linear relationship
between y and x .

The **linear equation** for predicting y^* from x^*
that minimizes mean squared error is simply:

$$\hat{y}_i^* = r_{xy} x_i^*$$

Thus, if x is observed to be 1 standard deviation above its own mean,
then we should predict that y will be r_{xy} standard deviations above its own mean.

$$\hat{y}_i^* = r_{xy} x_i^*$$

1 Means

2 Standard deviations

3 Correlation coefficient

$$\frac{(\hat{y}_i - \bar{y})}{\sigma_y} = r_{xy} \frac{(x_i - \bar{x})}{\sigma_x}$$

$$(\hat{y}_i - \bar{y}) = r_{xy} \frac{\sigma_y}{\sigma_x} (x_i - \bar{x})$$

$$(\hat{y}_i - \bar{y}) = r_{xy} \frac{\sigma_y}{\sigma_x} (x_i) - r_{xy} \frac{\sigma_y}{\sigma_x} (\bar{x})$$

$$\hat{y}_i = r_{xy} \frac{\sigma_y}{\sigma_x} (x_i) - r_{xy} \frac{\sigma_y}{\sigma_x} (\bar{x}) + \bar{y}$$

$$\hat{y}_i = \bar{y} - r_{xy} \frac{\sigma_y}{\sigma_x} (\bar{x}) + r_{xy} \frac{\sigma_y}{\sigma_x} (x_i)$$

$$\beta_1 = r_{xy} \frac{\sigma_y}{\sigma_x}$$

Where:

$$\beta_0 = \bar{y} - \beta_1 \bar{x}$$

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

$$\hat{y}_i = \bar{y} - \beta_1 \bar{x} + \beta_1 x_i$$

R Squared

Residual (Error) Sum of Squares **SSE**

Difference between actual and predicted values of y





$$R^2(y, \hat{y}) = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$




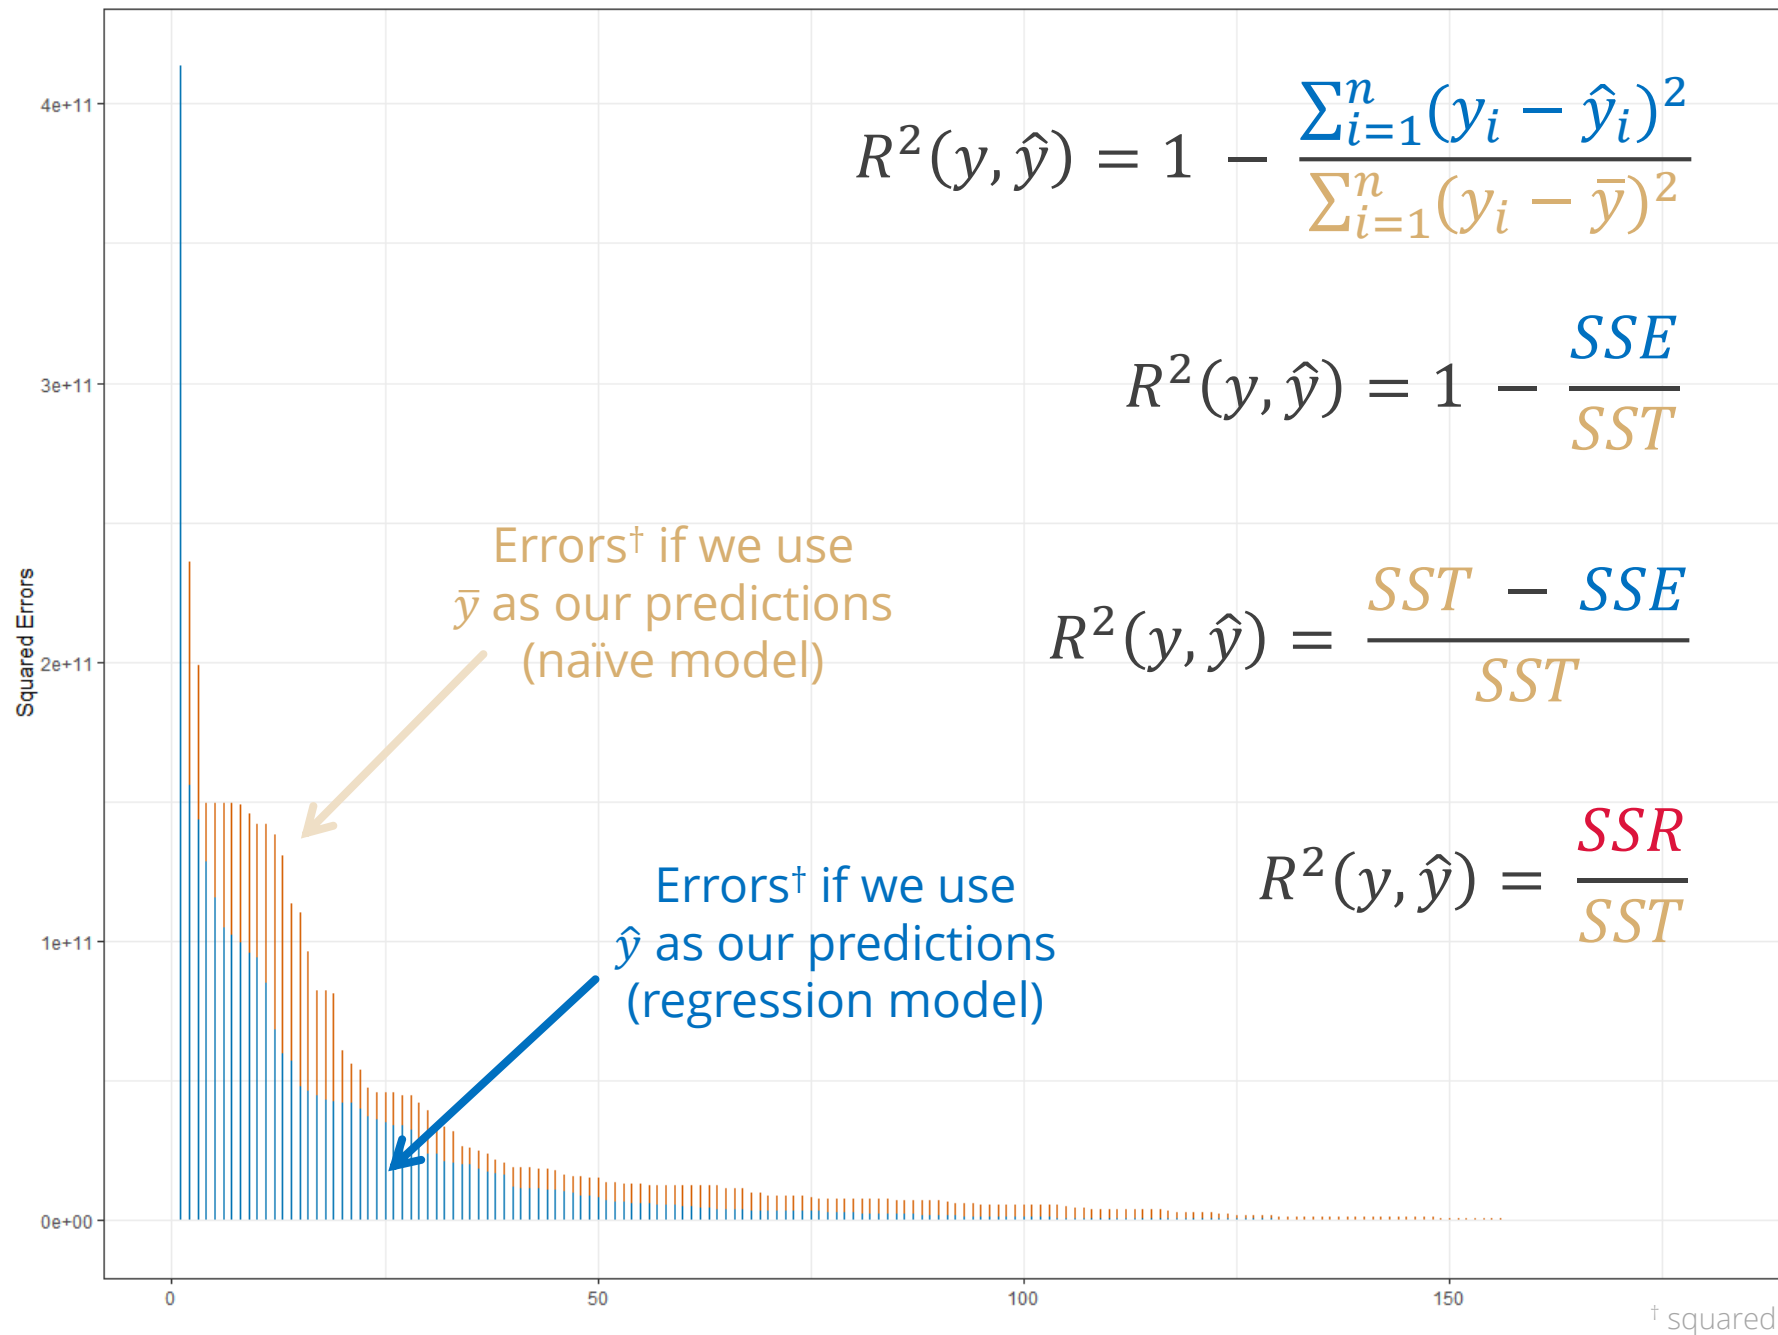
Fluctuations in y from its mean

Total Sum of Squares **SST**









$$R^2(y, \hat{y}) = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

$$R^2(y, \hat{y}) = 1 - \frac{SSE}{SST}$$

$$R^2(y, \hat{y}) = \frac{SST - SSE}{SST}$$

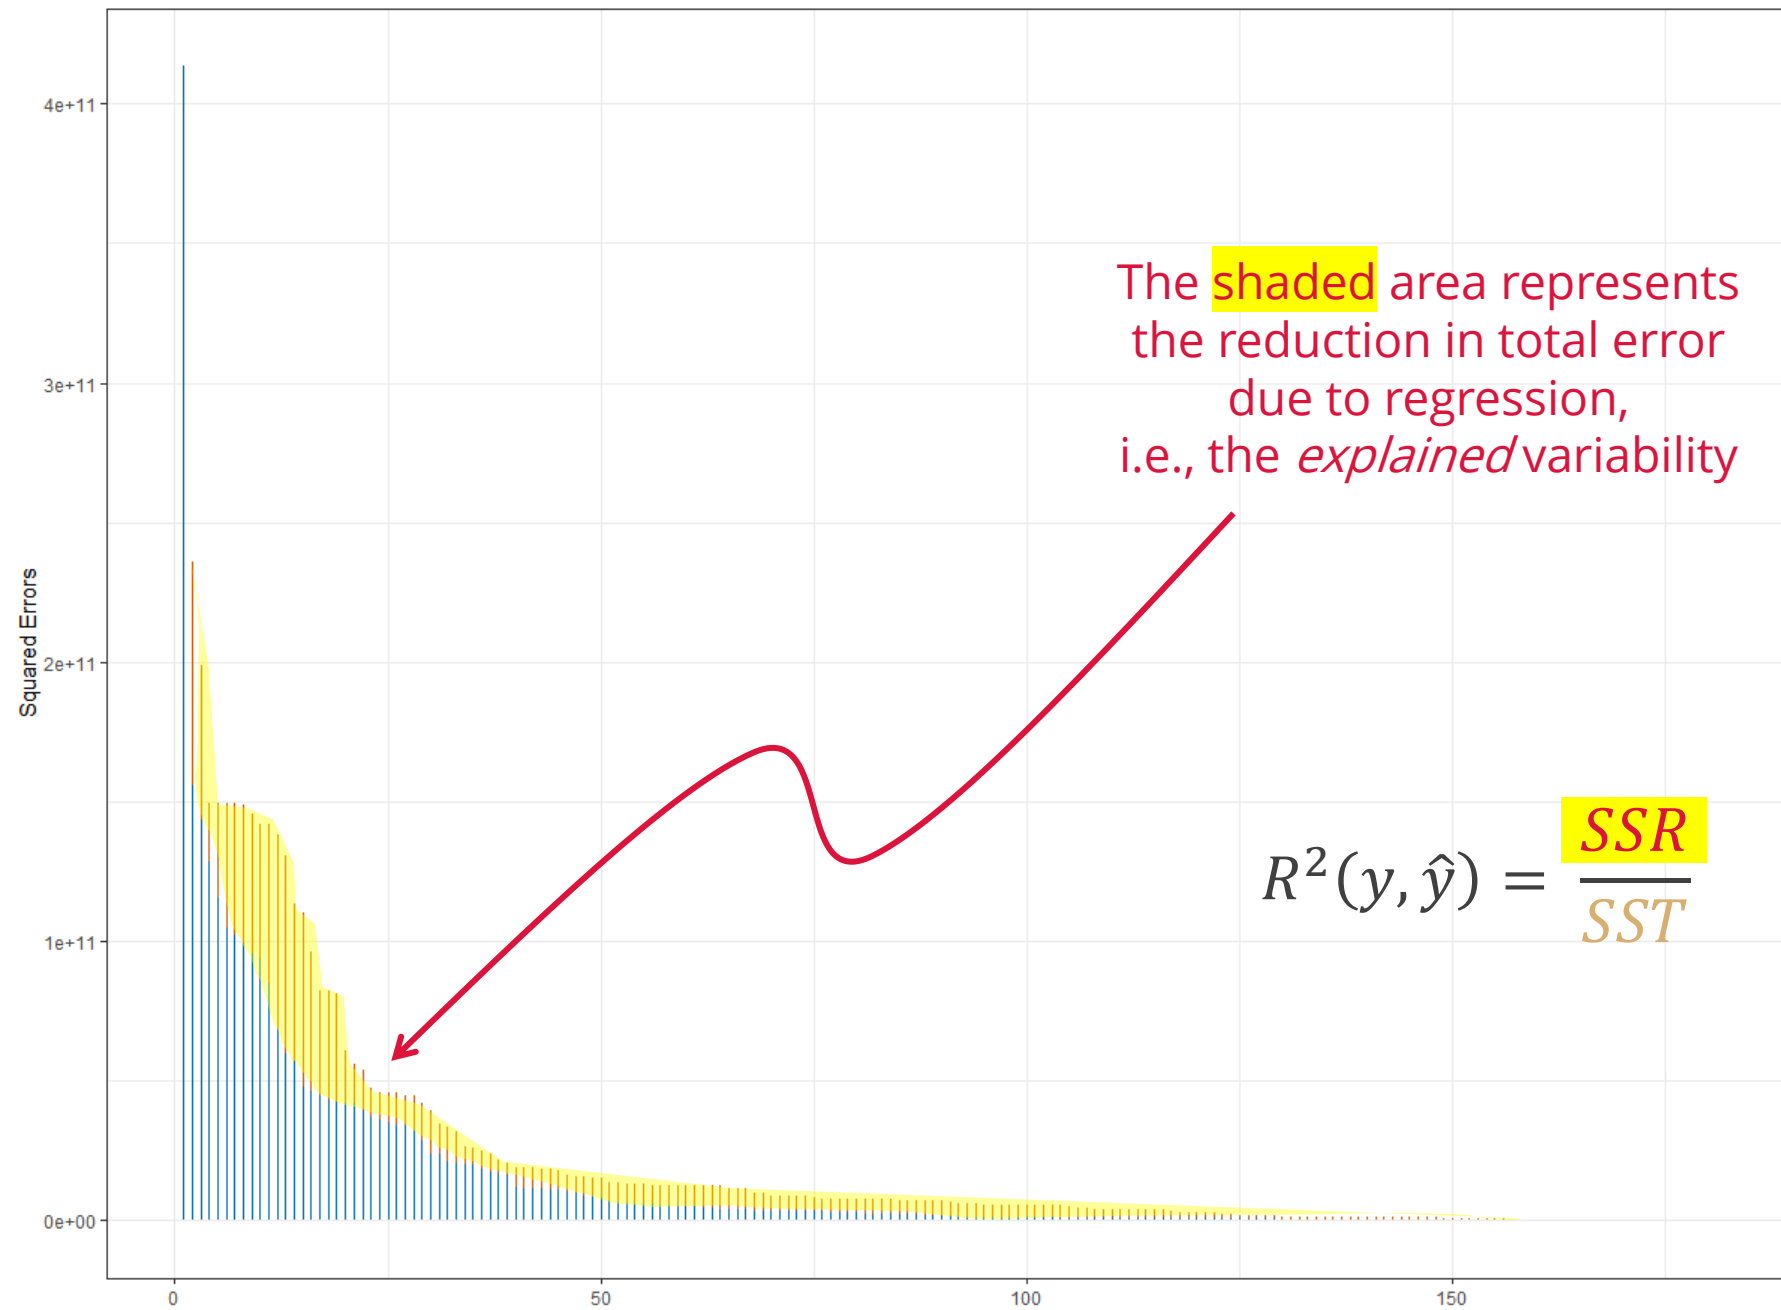
$$R^2(y, \hat{y}) = \frac{SSR}{SST}$$

Variability


SST = Total

SSR = Explained

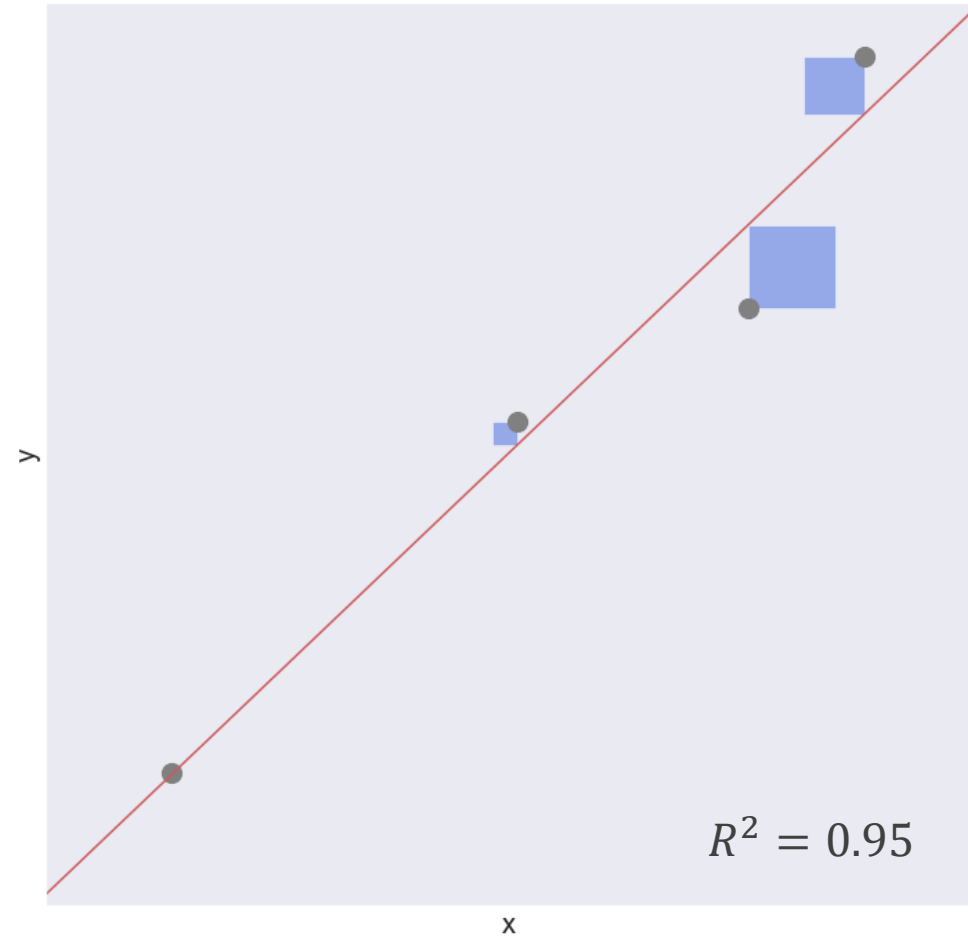
SSE = Unexplained



SST




SSE




$$R^2 = 1 - \frac{SSE}{SST}$$

$$R^2 = 0.95$$



Polynomial Regression ($n=4$)



Linear Regression

t Statistic

t-statistic

$$t = \frac{b_1 - \beta_1^{(0)}}{SE_{b_1}}$$

Student's *t*-distribution

Variety



vs.



Consistency




**Statistical
Quality Control**



William Sealy Gosset

A 19th century English statistician



The Probable Error of a Mean (1908), published by an anonymous "Student"

Quality of Hops → Quality of Beer

	Batch 1	Batch 2
Soft resins content	8.1% (n=11)	8.4% (n=14)

- **How many observations** were necessary to be confident, e.g., is the soft resins content of the extract within some small degrees of the desired level?
- **Is the difference statistically significant?**

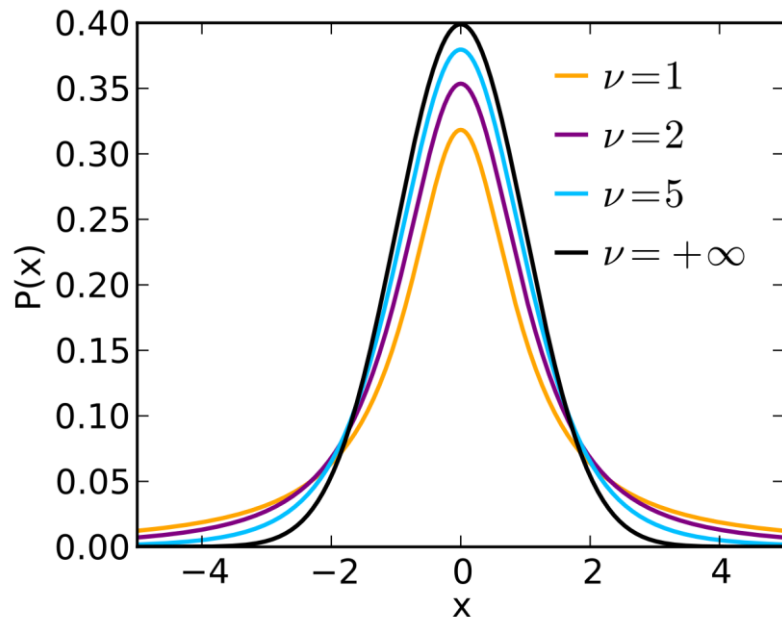


Student's t

t-statistic

$$t = \frac{\hat{\beta} - 0}{SE_{\hat{\beta}}}$$

It measures
how many **standard deviations** away from zero
the **estimated coefficient** is.




The ***p*-value** is the probability
of observing a t -statistic that large or larger in magnitude
given the null hypothesis that the true coefficient value is zero.


***p*-value > 0.05** → the variable is “accidentally” significant


Useful for Feature Selection.

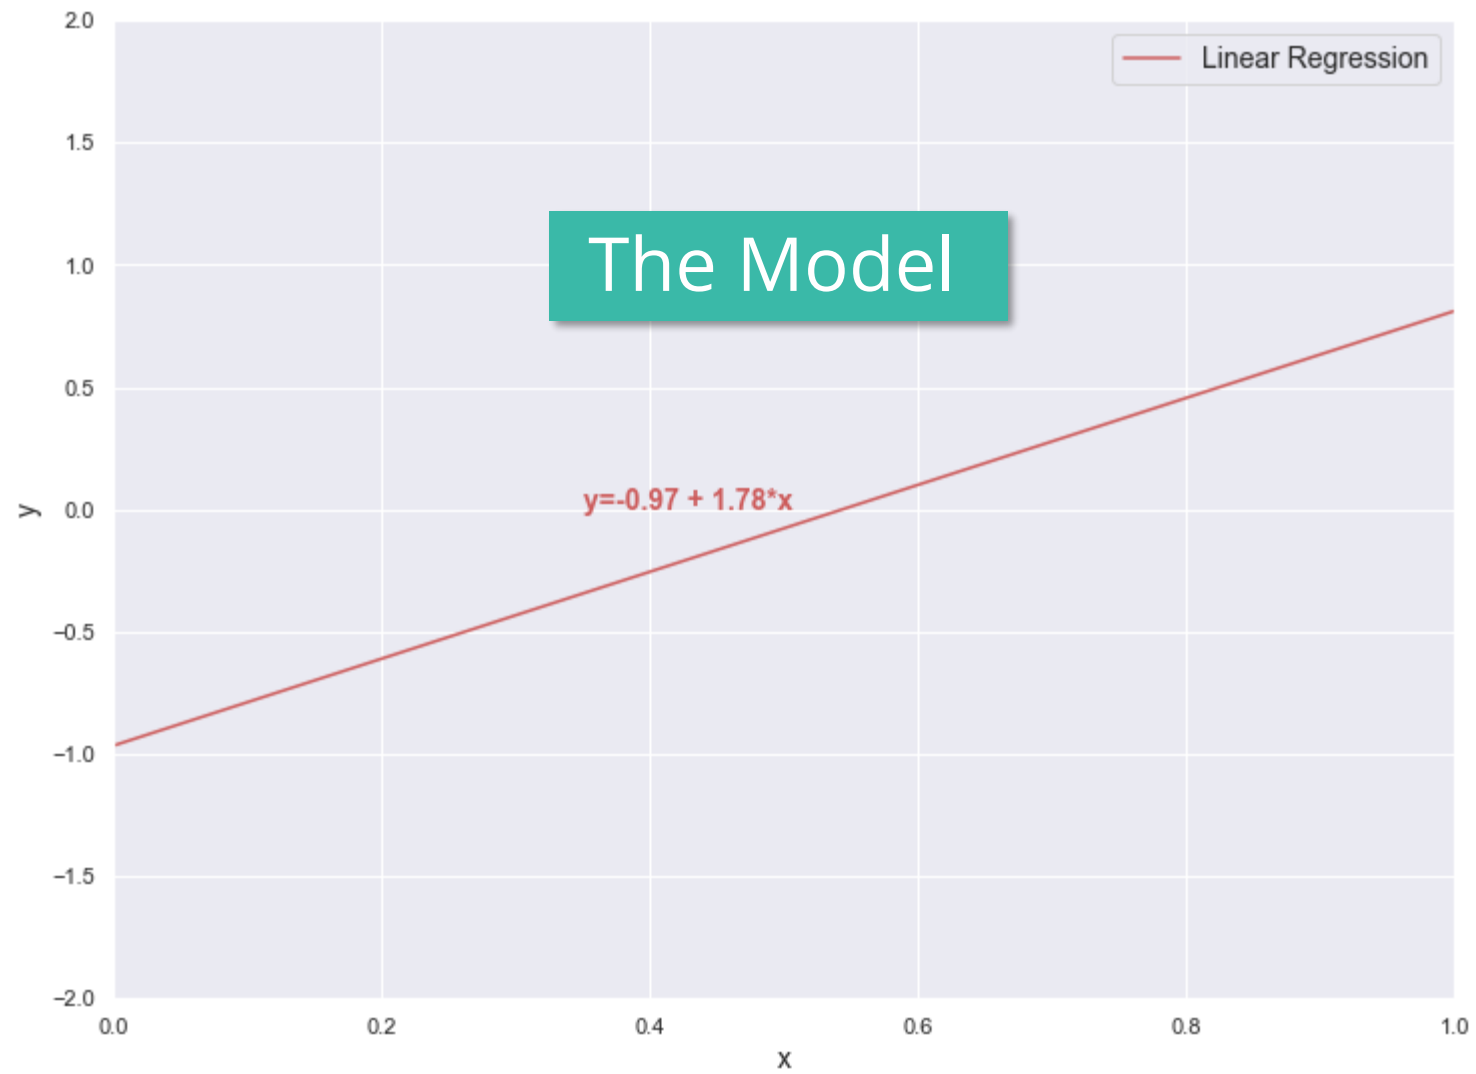
Simple Linear Regression (Example)





$$y = -\sin(2.2 * \pi * X) + \varepsilon$$

















MSE = **0.29**

R-Squared = **0.31**

Linear Regression

DATA SET

$$\{y_i, x_{i1}, \dots, x_{ij}\}_{i=1}^n$$

EQUATION

$$y = \underbrace{X^T \beta}_{\text{linear}} + \varepsilon$$

The model is linear in its parameters.

ASSUMPTION

$$\varepsilon \sim N(0, \sigma^2)$$

The error is a Gaussian random variable with expectation zero and variance σ^2 .

$$y = \begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ \vdots \\ \vdots \\ y_n \end{pmatrix} \quad X = \begin{pmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1j} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2j} \\ x_{31} & x_{32} & x_{33} & \dots & x_{3j} \\ \vdots & \vdots & \vdots & \dots & \vdots \\ \vdots & \vdots & \vdots & \dots & \vdots \\ \vdots & \vdots & \vdots & \dots & \vdots \\ x_{n1} & x_{n2} & x_{n3} & \dots & x_{nj} \end{pmatrix}$$

Supervised Learning

Linear Regression in `scikit-learn`

scikit-learn

Machine Learning in Python

Getting Started

What's New in 0.22.1

GitHub


- Simple and efficient tools for predictive data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more...




Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more...



Examples

Clustering

Automatic grouping of similar objects into sets.

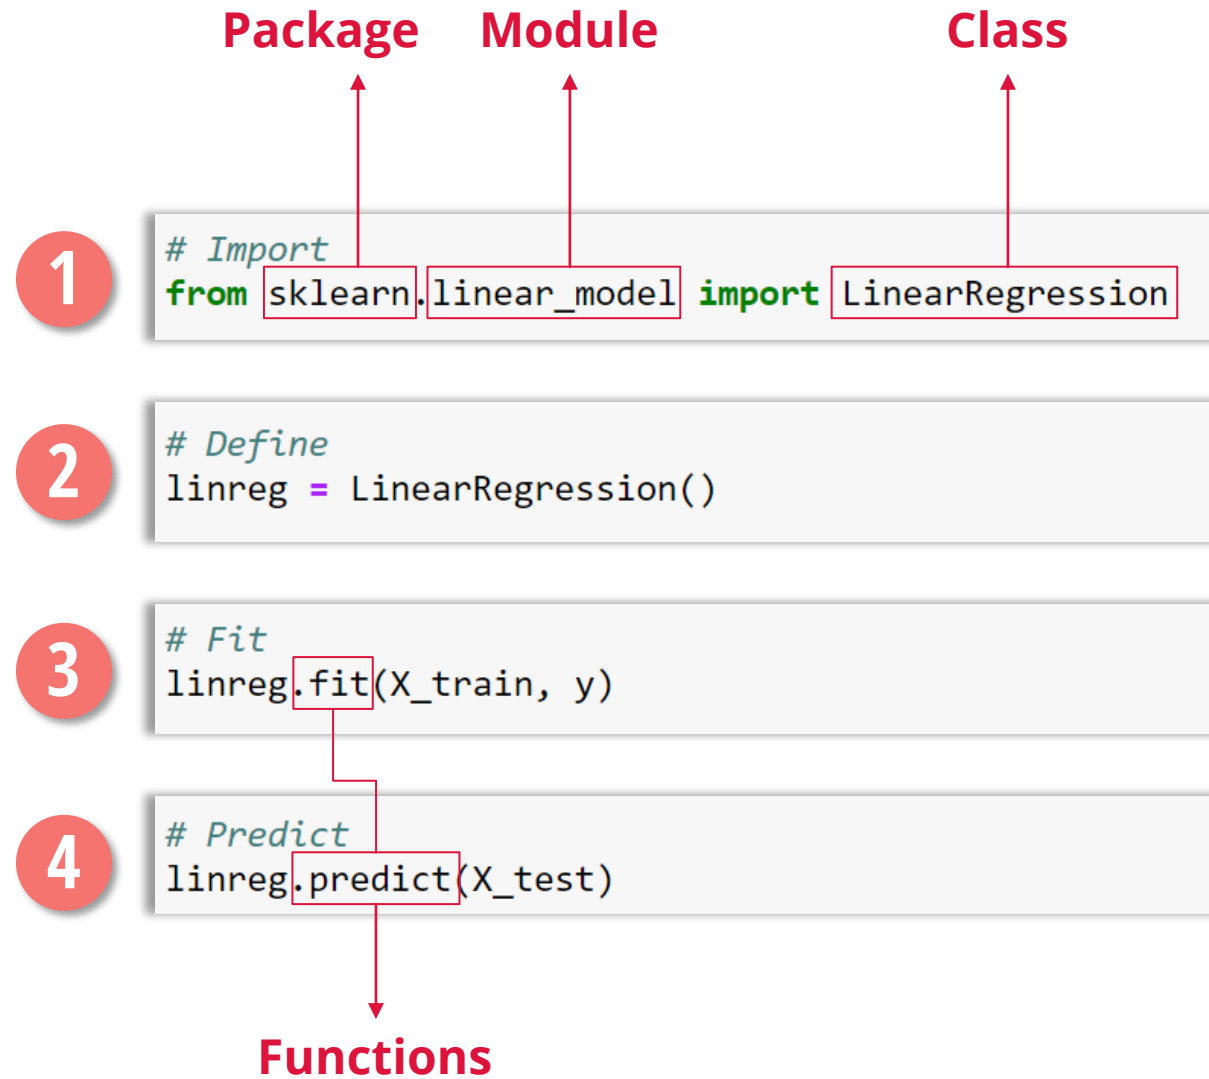
Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more...



Examples

```
class sklearn.linear_model.LinearRegression(  
    fit_intercept=True,  
    normalize=False,  
    copy_X=True,  
    n_jobs=None)
```




```
class sklearn.linear_model.LinearRegression(
```

```
    fit_intercept=True,
```

```
    normalize=False,
```

```
    copy_X=True,
```

```
    n_jobs=None)
```

Whether to calculate the intercept for this model.

If set to **False**,

no intercept will be used in calculations
(e.g. data is expected to be already centered)

Recommendation: `fit_intercept = True` (default)

```
class sklearn.linear_model.LinearRegression(  
    fit_intercept=True,  
    normalize=False,  
    copy_X=True,  
    n_jobs=None)
```

If **True**, the regressors X will be normalized before regression by subtracting the mean and dividing by the l2-norm.

This parameter is ignored when **fit_intercept** is set to **False**.

Recommendation: **normalize = False** (default)

Normalize the data prior to training a model.


Linear Regression Tutorial

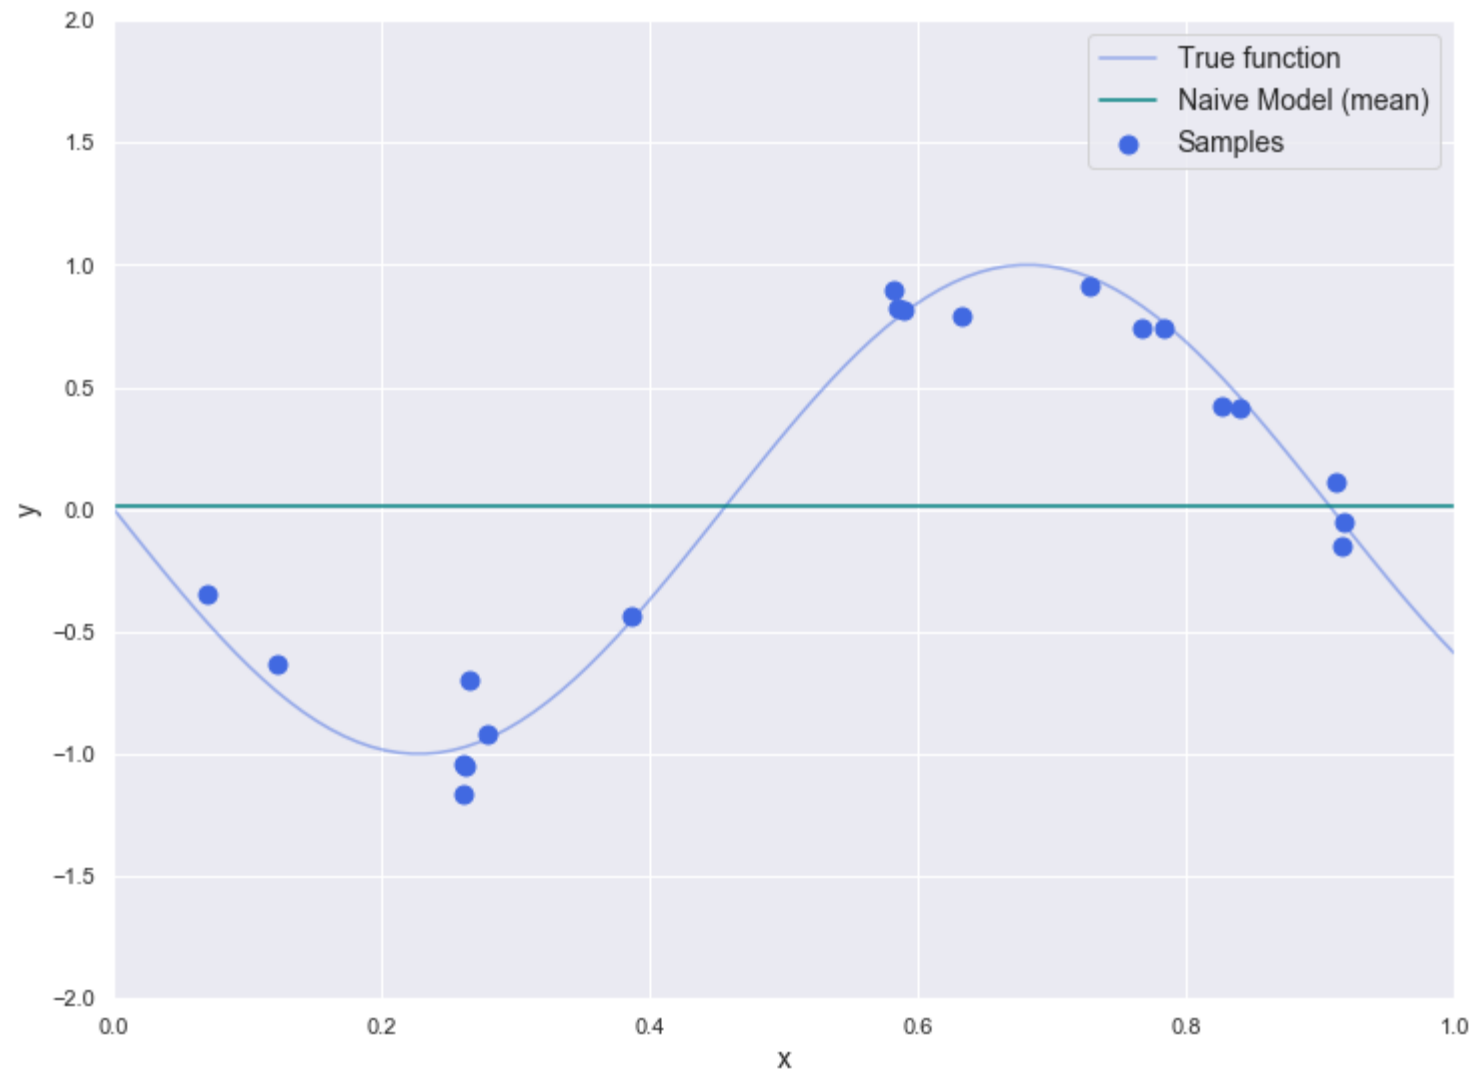
08_linear_reg_intro.ipynb



If she loves you more each and every day,
by linear regression she hated you before you met.

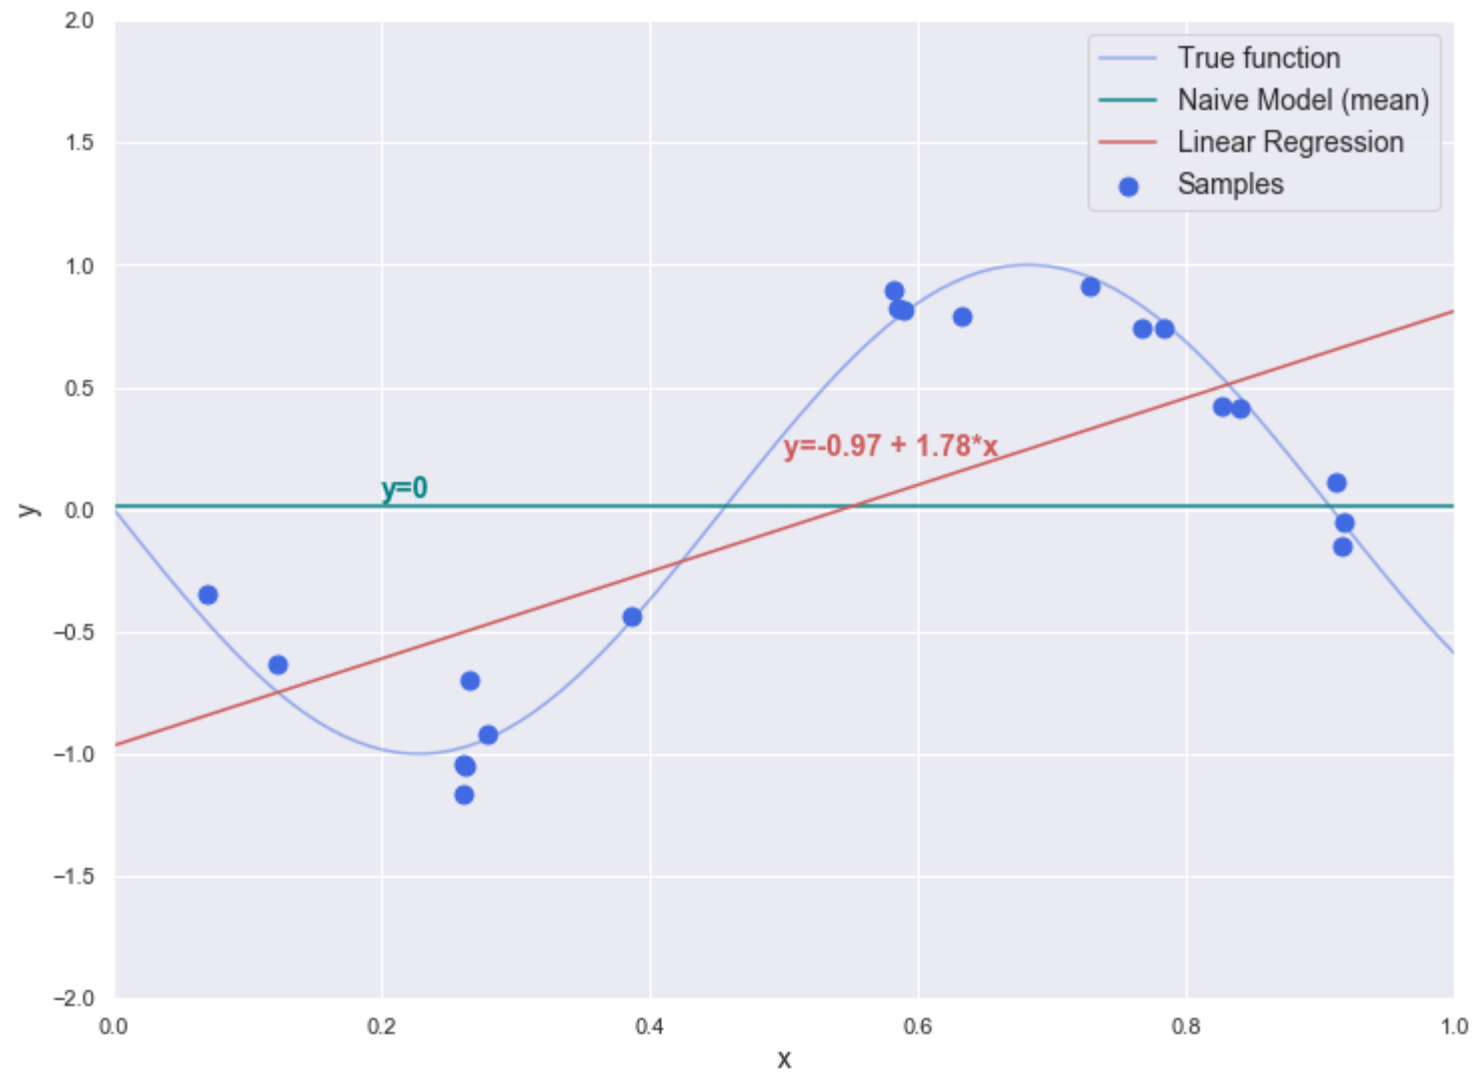
MY HOBBY: EXTRAPOLATING






MSE = **0.53**

R-Squared = **0.00**




MSE = 0.29


MSE = 0.53




MSE = **0.11**




MSE = **0.09**




MSE = **0.05**







MSE = **0.29**

MSE = **0.05**






Decision Tree Regressor



Recursive Partitioning



Depth = 0





$$x \leq 0.48$$

$$x > 0.48$$

Split rules



Depth = 1




$$x \leq 0.8$$

$$x > 0.8$$

Depth = 2





$J1, J2, J3$ = Leaf (terminal) nodes

1 How to partition the data?


2 When to stop?

1


How to partition the data?

2


When to stop?



Option 1





Option 2




...

Option n






Option 1



MSE = 0.15


Option 2




MSE = 0.11

...


Option n



MSE = 0.53




Option 1



MSE = 0.15


Option 2



MSE = 0.11

...

Option n



MSE = 0.53

LINEAR REGRESSION

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2$$

DECISION TREE REGRESSION

$$MSE = \frac{1}{n} \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

Decision Trees in `scikit-learn`


```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

1

```
# Import  
from sklearn.tree import DecisionTreeRegressor
```

2

```
# Define  
tree = DecisionTreeRegressor
```

3

```
# Fit  
tree.fit(X_train, y)
```

4

```
# Predict  
tree.predict(X_test)
```

```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

The function to measure the quality of a split.

'squared_error' = Mean Squared Error

$$MSE = \frac{1}{n} \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

1 How to partition the data?

2 When to stop?

```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

The maximum depth of the tree.

If **None**, then nodes are expanded until all leaves are pure or until all leaves contain less than `min_samples_split` samples.

Recommendation: `max_depth` start between 6 and 10

If `max_depth` is set to 3...


Depth = 0



Depth = 1



Depth = 2




Depth = 3



Depth = 4



Decision Tree Regression



```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

The minimum number of samples required to split an internal node:

If `int`, then consider `min_samples_split` as the minimum number.

If `float`, then $\text{ceil}(\text{min_samples_split} * \text{n_samples})$ are the minimum number of samples for each split.





Recommendation: `min_samples_split = 0.05`


```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

The minimum number of samples required to be at a leaf node.

A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches.

Recommendation: `min_samples_leaf = 0.02`



Should a split be considered?

If a split results in a children node with less than `min_samples_leaf` records, then discard that split.

```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

The number of features to consider when looking for the best split.

- If int, then consider max_features features at each split.
- If float, then max_features is a fraction
and $\text{int}(\text{max_features} * \text{n_features})$ features are considered at each split.
- If "auto", then max_features=n_features.
- If "sqrt", then max_features=sqrt(n_features).
- If "log2", then max_features=log2(n_features).
- If None, then max_features=n_features.

Recommendation: Consider using 'sqrt' or 'log2' if training on a large dataset; otherwise, leave to None.

```
class sklearn.tree.DecisionTreeRegressor(  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0)
```

**Set a user-defined seed
for reproducible results.**

If **int**, **random_state** is the seed used by
the random number generator.

Recommendation: Always set a seed (e.g., 314) to
ensure reproducible results.

Decision Tree Tutorial

05_decision_tree_intro.ipynb

Decision Tree Algorithm

PSEUDOCODE

1. Start at the root node.
2. For each feature:
 - Identify the best split that minimizes MSE.
3. Identify the feature that generates the lowest MSE.
4. Split the node using that feature and its best split.
5. Repeat steps 2 thru 4 until a stopping criterion is met.

ID3

Iterative Dichotomizer



C4.5



CART

Classification and
Regression Trees

sklearn

Decision Trees

- Simple and intuitive
- Can handle non-linear relationships
- Can handle missing values
- Can handle both numeric and categorical variables[†]
- Not influenced heavily by outliers

[†] However, in scikit-learn, you still need to convert them into a numeric form.

However...

○ **Decision tree** is a **greedy** algorithm; it tends to **overfit** on data with large number of features.

Recommendations:

1. Perform **feature reduction** before training a model.
2. Always use `min_samples_leaf` to control the amount of over-fitting.
3. Try a small tree first, using `max_depth`, and then grow further if necessary.

Summary

1. Regression to mediocrity
2. Linear Regression model
 - Correlation coefficient
 - R^2
 - Sum of squares
 - t statistic
3. Linear Regression in `scikit-learn`
4. Build a model using partitioning
5. Decision trees
6. Decision trees in `scikit-learn`

Data Science / ML Term	Multidisciplinary Synonyms
Label	<ul style="list-style-type: none"> • Dependent variable • Response variable • Target • Output
Features	<ul style="list-style-type: none"> • Independent variables • Explanatory variables • Attributes • Inputs • Predictors
(Regression) Coefficients	<ul style="list-style-type: none"> • Parameter estimates • Slopes
Noise	<ul style="list-style-type: none"> • Random error • Residuals
Cases	<ul style="list-style-type: none"> • Observations • Records • Rows
Train	<ul style="list-style-type: none"> • Fit • Build

Next Up

1. Introduction

2. The Data Science Process

3. Supervised Learning: Classification

4. Unsupervised Learning

5. The Grunt Work

6. Wrap Up