

# Supervised Learning: Classification

Spring 2023

1. Introduction

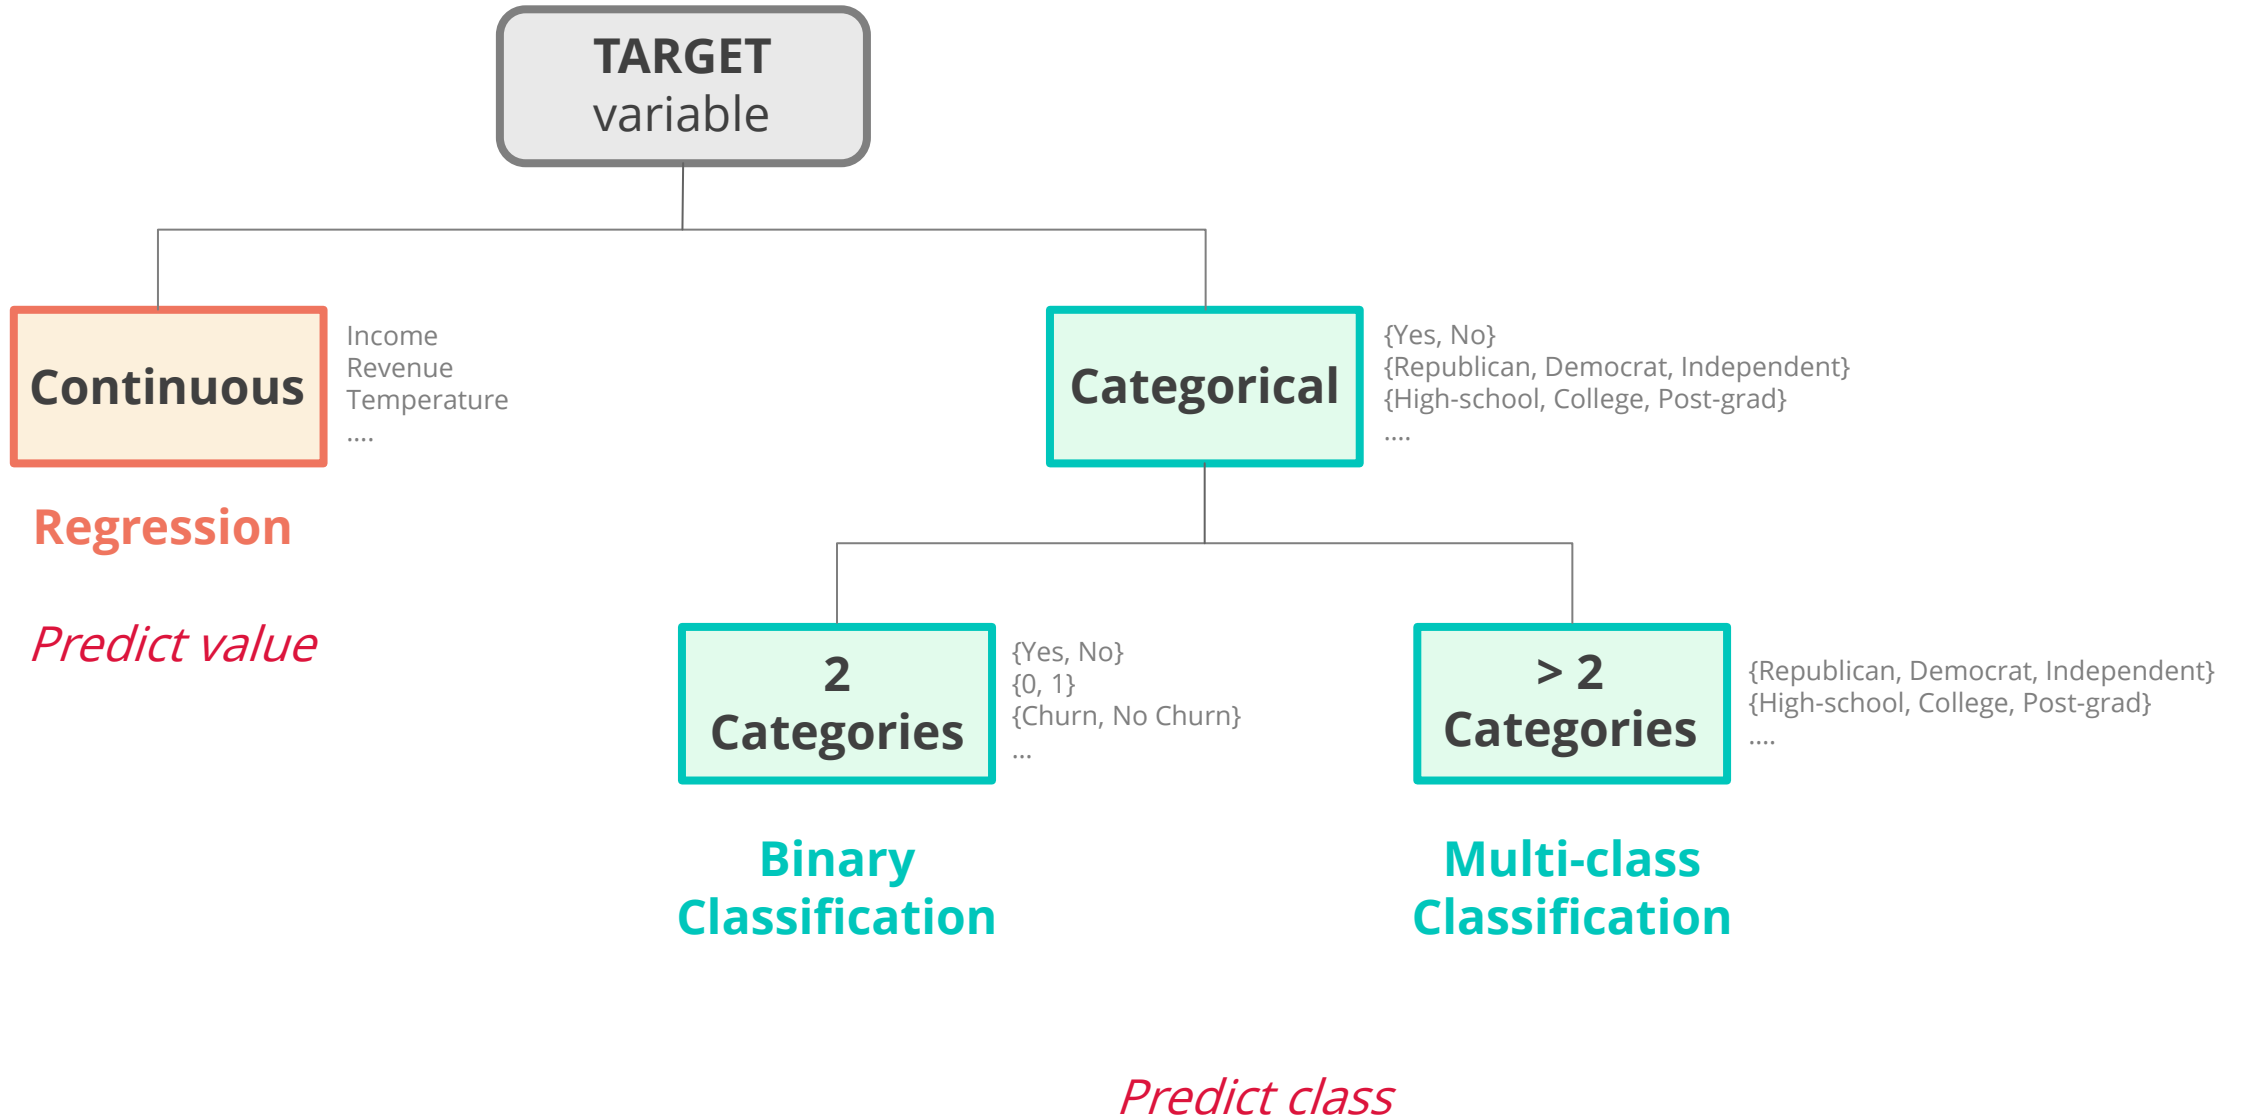
2. The Data Science Process

**3. Supervised Learning: Classification**


4. Unsupervised Learning

5. The Grunt Work


6. Wrap Up




# Classification Trees



Recursive partitioning on continuous data (outcome)



Binary outcome ( $y$ )



Recursive partitioning

1

How to partition the data?

2

When to stop?

## REGRESSION TREES


$$MSE = \frac{1}{n} \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

## CLASSIFICATION TREES


?

Measure of Impurity







Binary outcome ( $y$ )



Recursive partitioning




**Split #1**




**Split #2**


How to determine which split is better?



$$p_{y=0} = \frac{25}{100} = \mathbf{0.25}$$



$$p_{y=0} = \frac{50}{100} = \mathbf{0.50}$$




$$p_{y=0} = \frac{75}{100} = \mathbf{0.75}$$

$$p_{y=1} = \frac{75}{100} = \mathbf{0.75}$$

$$p_{y=1} = \frac{50}{100} = \mathbf{0.50}$$


$$p_{y=1} = \frac{25}{100} = \mathbf{0.25}$$

$p_{y=1}$  is not symmetrical.




$$p_{y=0} = \frac{25}{100} = \mathbf{0.25}$$

$$p_{y=1} = \frac{75}{100} = \mathbf{0.75}$$



$$p_{y=0} = \frac{50}{100} = \mathbf{0.50}$$

$$p_{y=1} = \frac{50}{100} = \mathbf{0.50}$$



$$p_{y=0} = \frac{75}{100} = \mathbf{0.75}$$

$$p_{y=1} = \frac{25}{100} = \mathbf{0.25}$$

**0.375**


**0.50**


**0.375**

$$p_{y=1} * (1 - p_{y=1}) + p_{y=0} * (1 - p_{y=0})$$


$$Entropy = - \sum_k p_k \log_2 p_k$$

$$Gini = \sum_k p_k (1 - p_k)$$






Where is the **optimal cut-off** (partition)  
that would yield the **lowest impurity**?

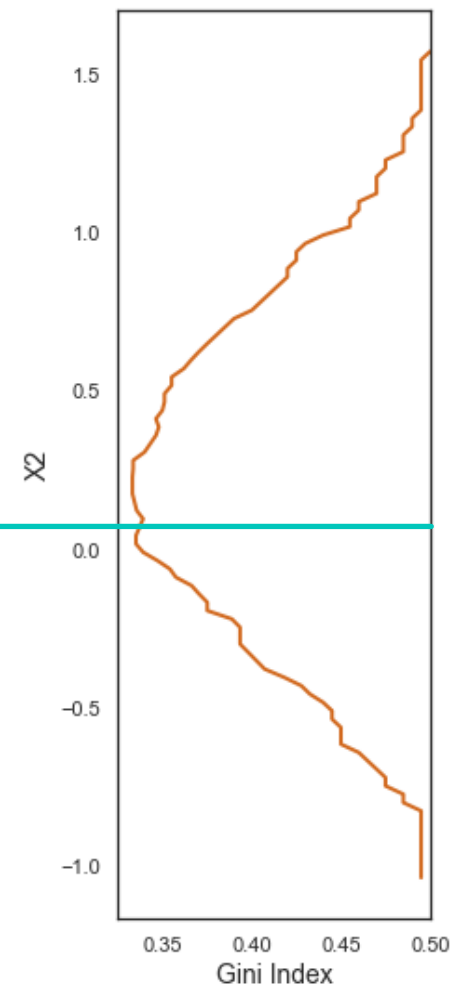



Min Gini = **0.33**




Min Gini = **0.36**







Min Gini = **0.36**







Min Gini = **0.33**



Min Gini = **0.33**

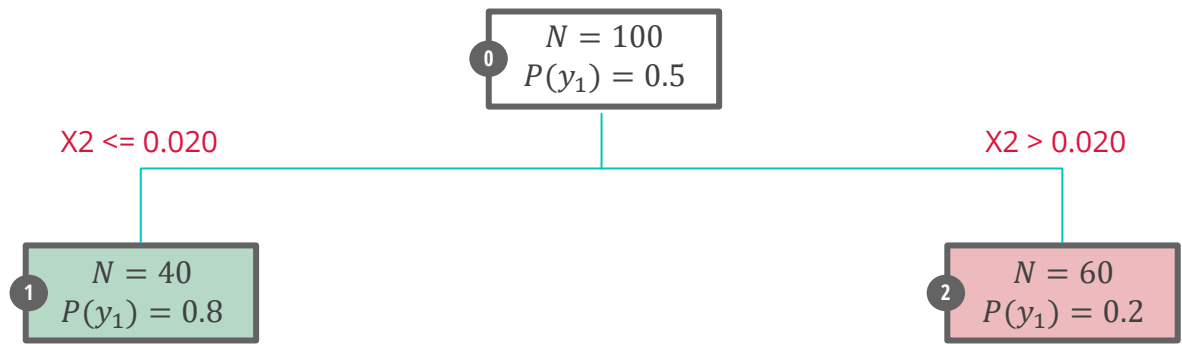

Min Gini = **0.36**

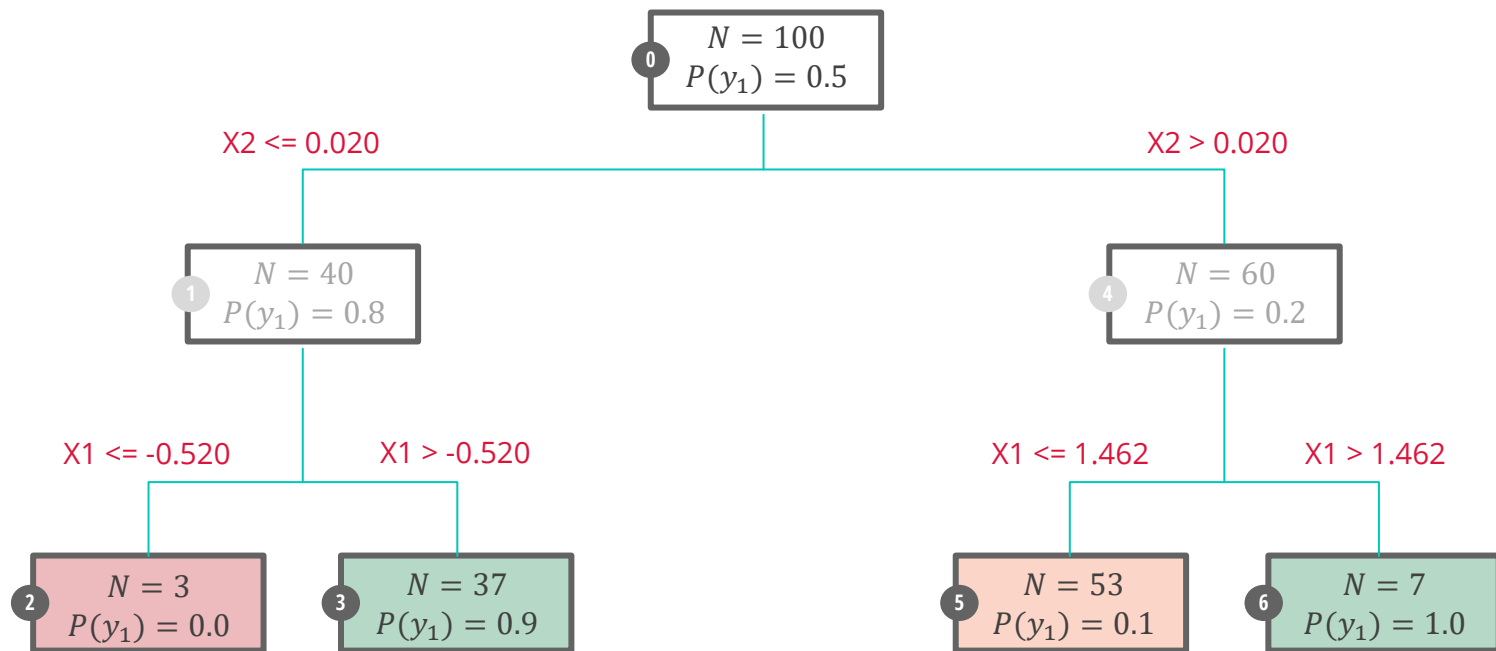



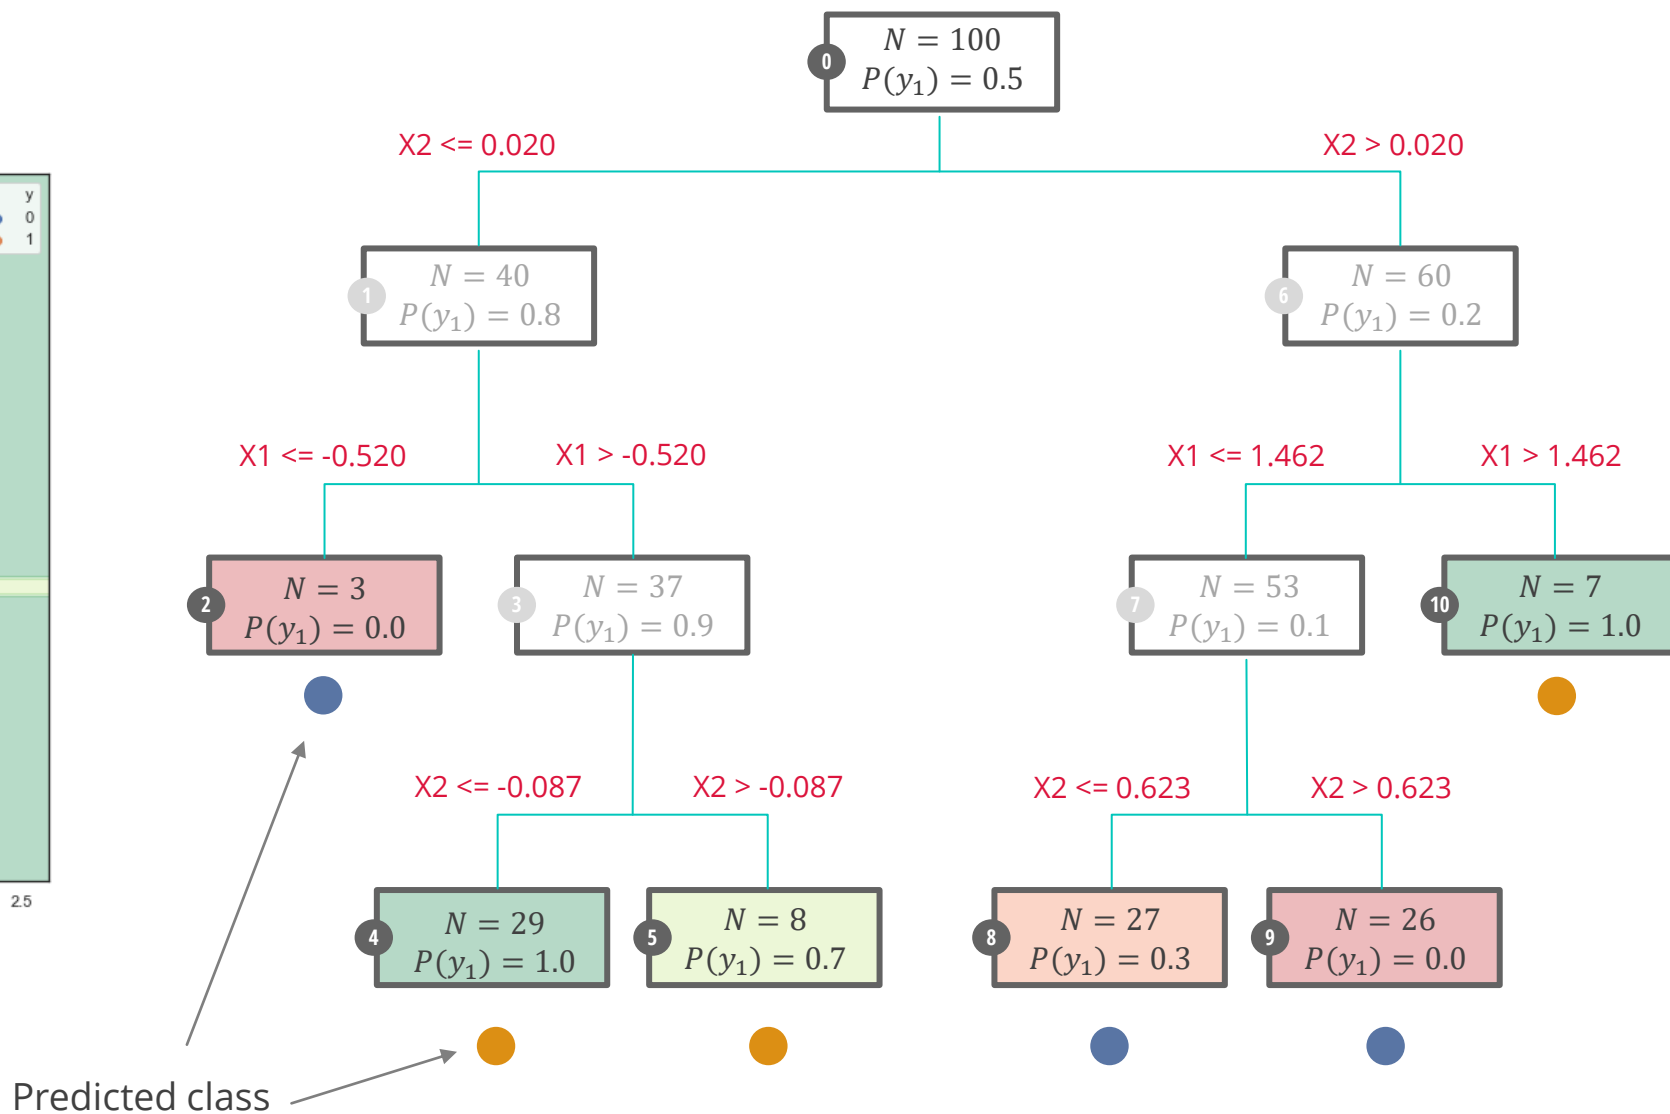



Reduction in impurity → “Information Gain”

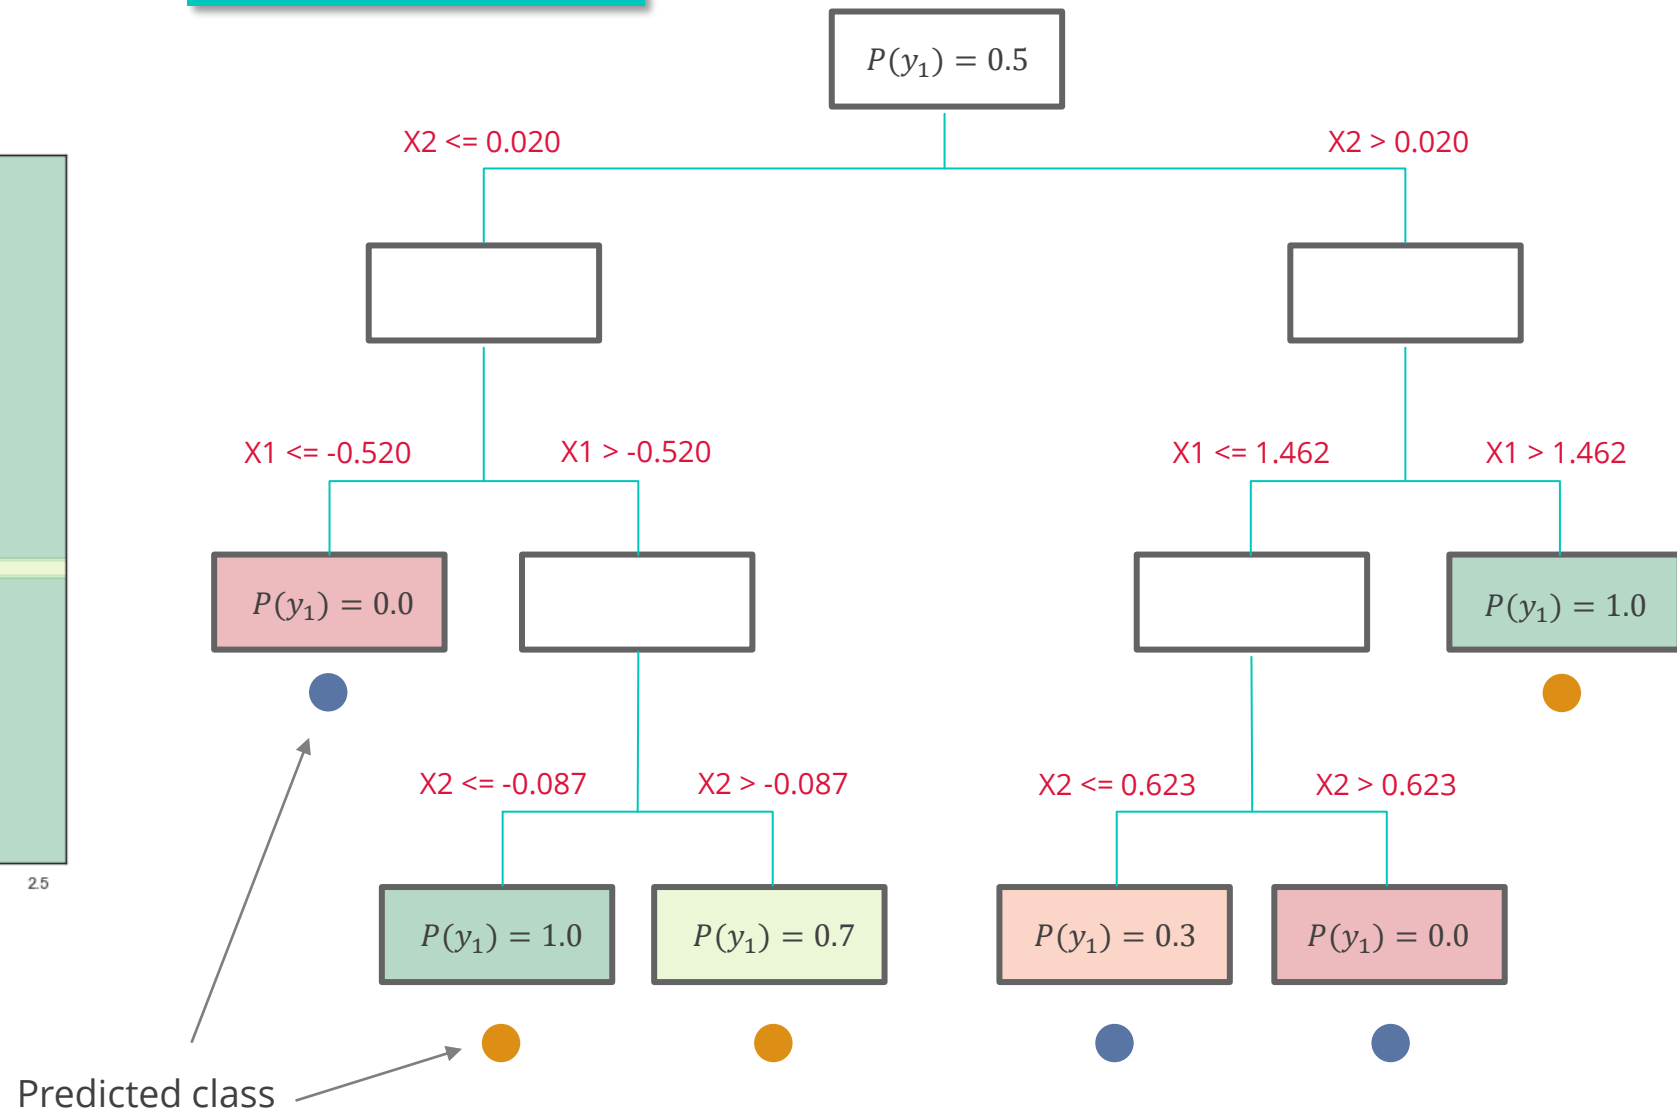



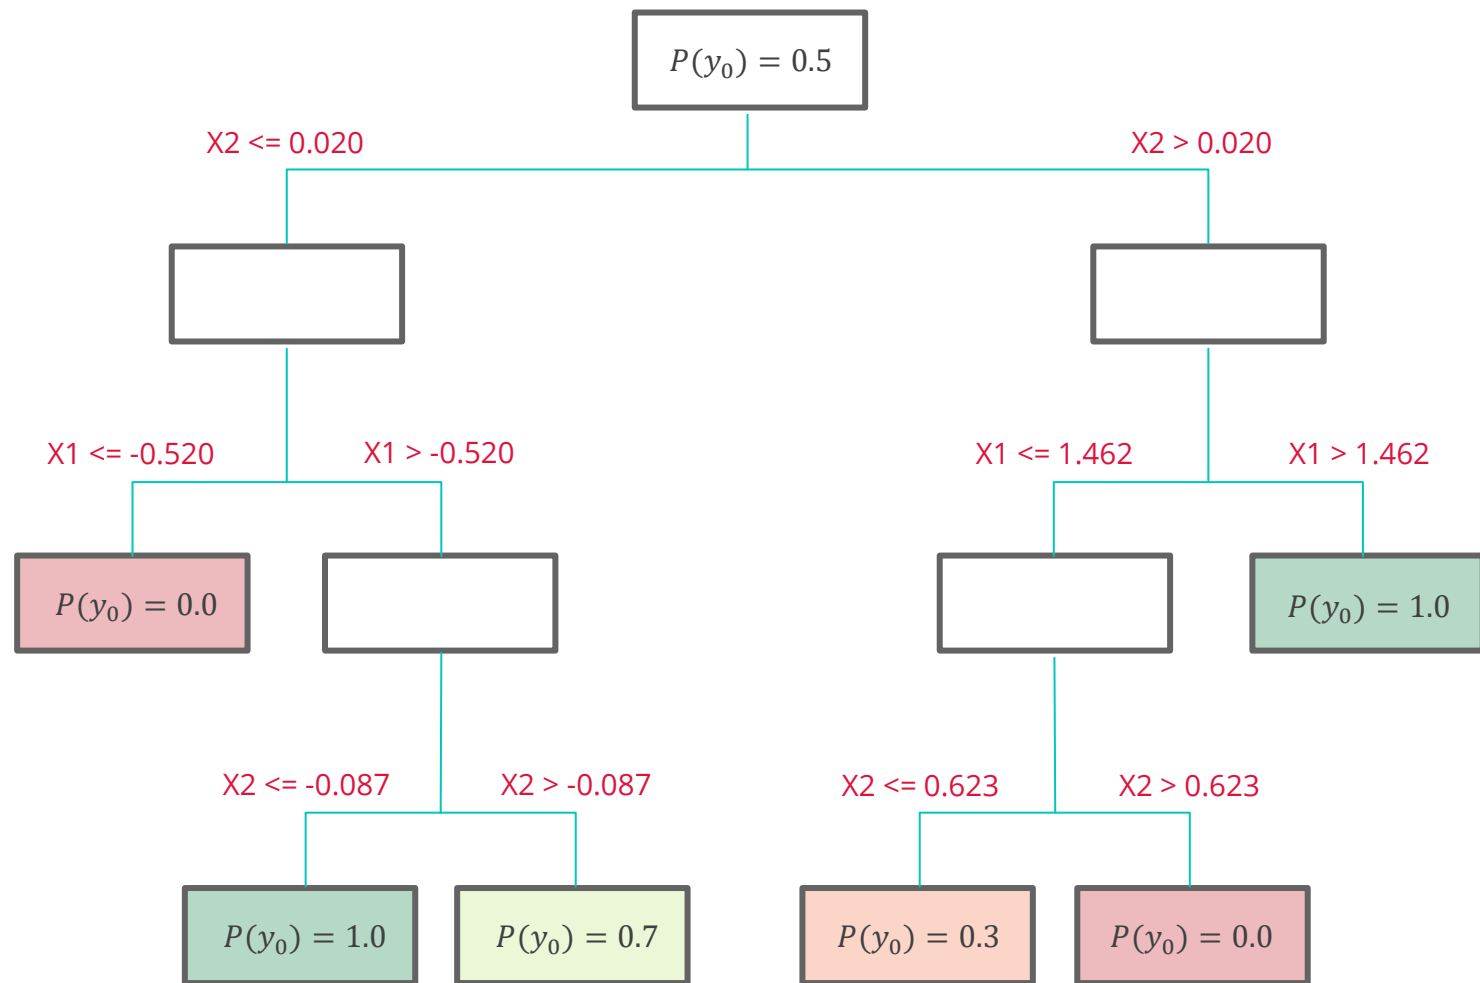







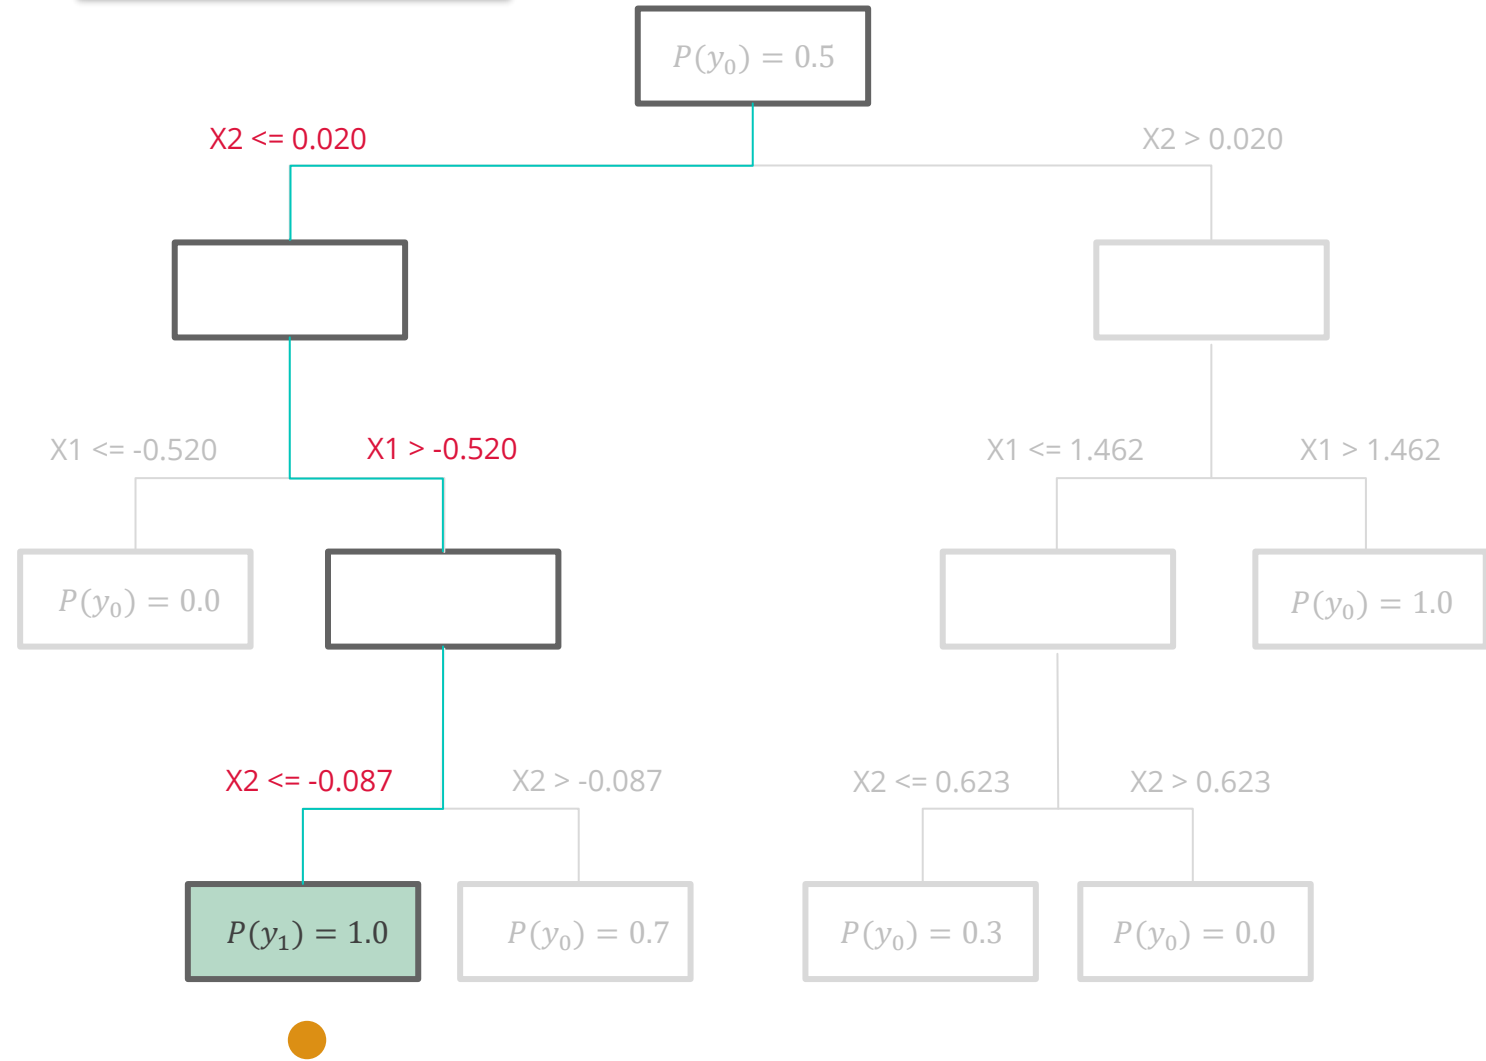

# The Model







# Prediction



# Decision Trees for Classification: Summary

Impurity measure	Formula
Entropy (Cross-entropy, deviance)	$-\sum_k p_k \log_2 p_k$
Gini index	$\sum_k p_k (1 - p_k)$
Mis-classification error	$\sum_k (1 - p_k)$
Chi-Square	$\sum_i \frac{(x_i - m_i)^2}{m_i}$


- **CHAID** (Chi-Square Automatic Interaction Detection), based on adjusted significant testing
  - Can perform two or more splits
  - Not available in `scikit-learn`

## Entropy

$$H = - \sum_k p_k \log_2 p_k$$

### Low Entropy


One class is more likely than the other.




**High purity**

### High Entropy


Both classes are nearly equally likely.



**High impurity**



$$p_{\square} = \frac{1}{7} \quad p_{\circ} = \frac{6}{7}$$



$$p_{\square} = \frac{2}{5} \quad p_{\circ} = \frac{3}{5}$$

$$H = - \sum_k p_k \log_2 p_k$$

$$H = - p_{\square} \log_2 p_{\square} - p_{\circ} \log_2 p_{\circ}$$

Low Entropy

$$H_A = 0.59$$


High purity

$$H_B = 0.97$$

High Entropy

Low purity

# Information Gain




$$IG = H - (H_A * P_A + H_B * P_B)$$

$$IG = 0.81 - \left( 0.59 * \frac{7}{12} + 0.97 * \frac{5}{12} \right)$$

$$IG = 0.81 - 0.75 = \mathbf{0.06}$$

- **Information Gain:** The amount by which the ambiguity (entropy) decreases due to the split.
- The goal is to find the split that **maximizes** the information gain.

# Information Entropy




**Claude Shannon**

(1916 – 2001)


The father of Information Theory

$$H = - \sum_k p_k \log_2 p_k$$



Information Entropy  
(7-minute [video](#))

1



A B A C B D B D C


$$P_A = 0.25$$

$$P_B = 0.25$$

$$P_C = 0.25$$

$$P_D = 0.25$$

2



D A C A B A A C A

$$P_A = 0.50$$

$$P_B = 0.125$$

$$P_C = 0.125$$

$$P_D = 0.25$$

Which machine is producing more **information**?


A

B

C

D

1



?

Which machine is  
producing more **information**?



2




?

If you had to predict the **next symbol** from  
a machine, **how many (yes/no) questions**  
you would have to ask?



1



?

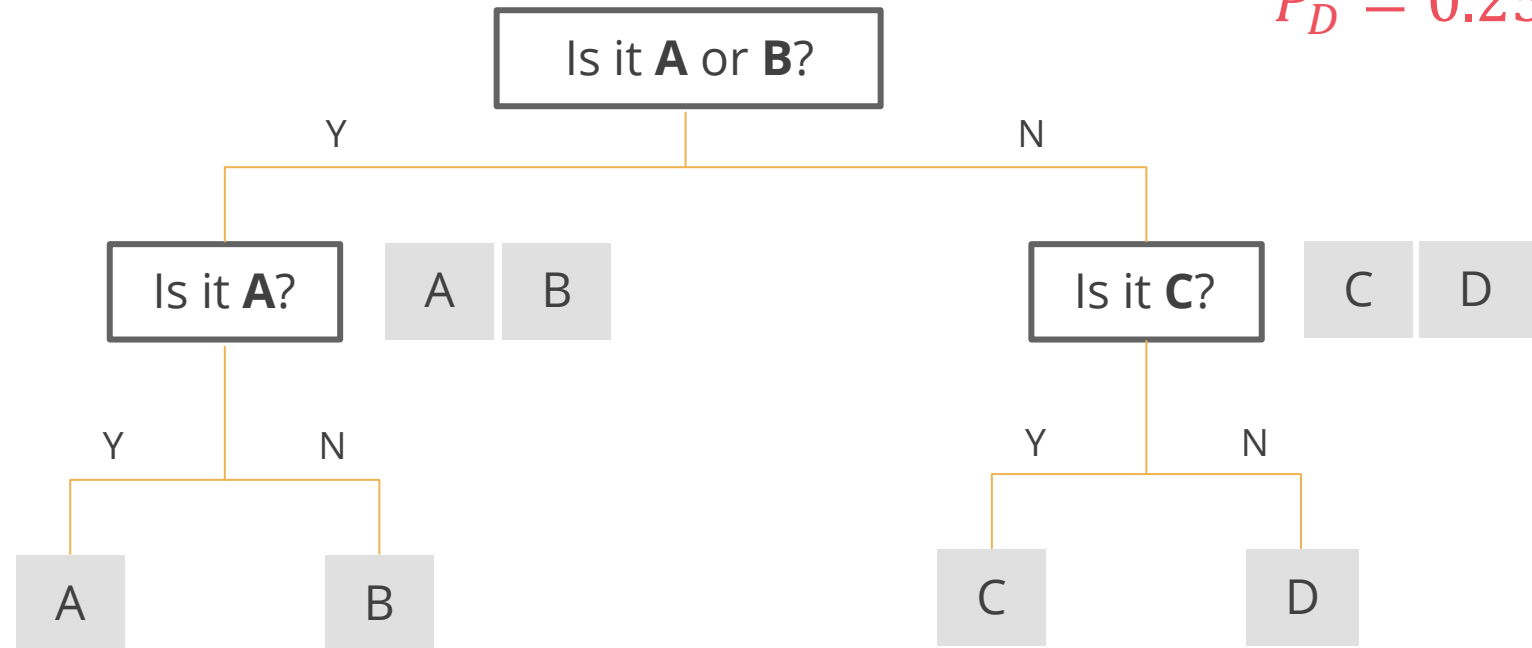
A B C D

$$P_A = 0.25$$

$$P_B = 0.25$$

$$P_C = 0.25$$

$$P_D = 0.25$$



The uncertainty in **machine 1** is **two questions** per symbol.

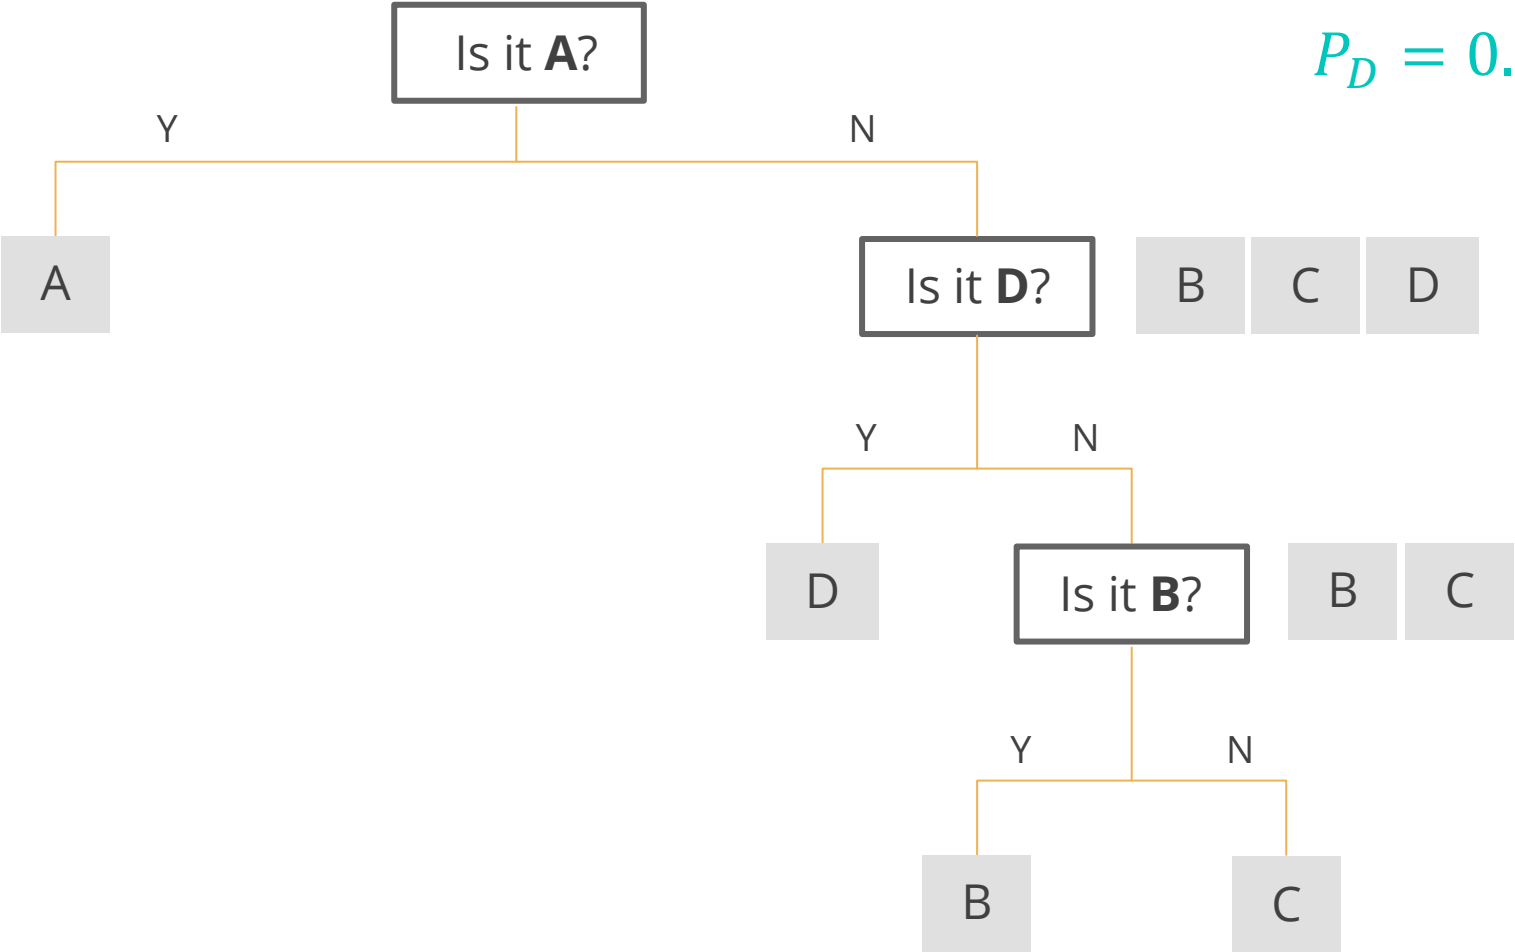
A B C D

$P_A = 0.50$


$P_B = 0.125$

$P_C = 0.125$

$P_D = 0.25$




2



?


A B C D

1




2 questions

2



?




$$P_A = 0.25$$

$$P_B = 0.25$$

$$P_C = 0.25$$

$$P_D = 0.25$$



$$P_A = 0.50$$

$$P_B = 0.125$$

$$P_C = 0.125$$


$$P_D = 0.25$$

$$\begin{aligned}
 \text{Expected \# of bounces} &= P_A * 1 + P_B * 3 + P_C * 3 + P_D * 2 \\
 &= 0.5 * 1 + 0.125 * 3 + 0.125 * 3 + 0.25 * 2 \\
 &= 1.75 \\
 &= \text{Expected \# of questions}
 \end{aligned}$$

2



?



$$P_A = 0.50$$

$$P_B = 0.125$$

$$P_C = 0.125$$

$$P_D = 0.25$$


A

B

C


D

1



2 questions

2



1.75 questions

**Machine 2** is producing **less information**, because there's less uncertainty (or surprise) about its output.

# A Mathematical Theory of Communication

By C. E. SHANNON





Fig. 7—Entropy in the case of two possibilities with probabilities  $p$  and  $(1 - p)$ .

$$H = \sum_k p_k * \text{\textit{\# of bounces}}_k$$



$$\text{\textit{\# of bounces}} = \log_2(\text{\textit{\# of outcomes}})$$

$$\text{\textit{\# of outcomes}} = \frac{1}{p}$$

$$\text{\textit{\# of bounces}} = \log_2\left(\frac{1}{p}\right) = -\log_2 p$$

$$H = - \sum_k p_k \log_2 p_k$$



# Classification Trees in `scikit-learn`



```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    ccp_alpha=0.0)
```

1

```
# Import  
from sklearn.tree import DecisionTreeClassifier
```

2

```
# Define  
clf = DecisionTreeClassifier()
```

3

```
# Fit  
clf.fit(x_train, y_class)
```

4

```
# Predict  
clf.predict(x_test)
```

```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    ccp_alpha=0.0)
```

**The function to measure  
the quality of a split.**

Supported criteria are “**gini**” for the Gini impurity and “**entropy**” for the information gain.

```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    ccp_alpha=0.0)
```

### The maximum depth of the tree.

If **None**, then nodes are expanded until all leaves are pure or until all leaves contain less than `min_samples_split` samples.

Recommendation: `max_depth` between 6 and 10

```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    ccp_alpha=0.0)
```

**The minimum number of samples required  
to split an internal node:**

If `int`, then consider `min_samples_split`  
as the minimum number.

If `float`, then  
`ceil(min_samples_split * n_samples)`  
are the minimum number of  
samples for each split.

Recommendation: `min_samples_split = 0.05`

```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    presort=False)
```

**The minimum number of samples required to be at a leaf node.**

A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches.

Recommendation: `min_samples_leaf = 0.02`

```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    ccp_alpha=0.0)
```

**Set a user-defined seed for  
reproducible results.**

If `int`, `random_state` is the seed used  
by the random number generator.


Recommendation: Always set a seed (e.g., 314) to  
ensure reproducible results.

```
class sklearn.tree.DecisionTreeClassifier(  
    criterion='gini',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    class_weight=None,  
    ccp_alpha=0.0)
```

**Weights associated with classes in the form {`class_label`: weight}.**

If not given, all classes are supposed to have weight one.

Recommendation: `class_weight = 'balanced'`  
*when class imbalance is high*







**Decision Tree  
is too small**

**Underfitting**

**High bias**



**BOOSTING**

**Decision Tree  
is too large**

**Overfitting**

**High variance**



**BAGGING**

# Random Forests

Distribution of the estimates of the dressed weight of a particular living ox, made by 787 different persons.


Degrees of the length of Array 0°—100°	Estimates in lbs.	Centiles		Excess of Observed over Normal
		Observed deviates from 1207 lbs.	Normal p.e = 37	
5	1074	-133	-90	+43
10	1109	-98	-70	+28
15	1126	-81	-57	+24
20	1148	-59	-46	+13
$q_1$ 25	1162	-45	-37	+8
30	1174	-33	-29	+4
35	1181	-26	-21	+5
40	1188	-19	-14	+5
45	1197	-10	-7	+3
$m$ 50	1207	0	0	0
55	1214	+7	+7	0
60	1219	+12	+14	-2
65	1225	+18	+21	-3
70	1230	+23	+29	-6
$q_3$ 75	1236	+29	+37	-8
80	1243	+36	+46	-10
85	1254	+47	+57	-10
90	1267	+52	+70	-18
95	1293	+86	+90	-4

$q_1$ ,  $q_3$ , the first and third quartiles, stand at 25° and 75° respectively.

$m$ , the median or middlemost value, stands at 50°.

The dressed weight proved to be 1198 lbs.

Diagram, from the tabular values.




The continuous line is the normal curve with p.e. = 37.

The broken line is drawn from the observations.

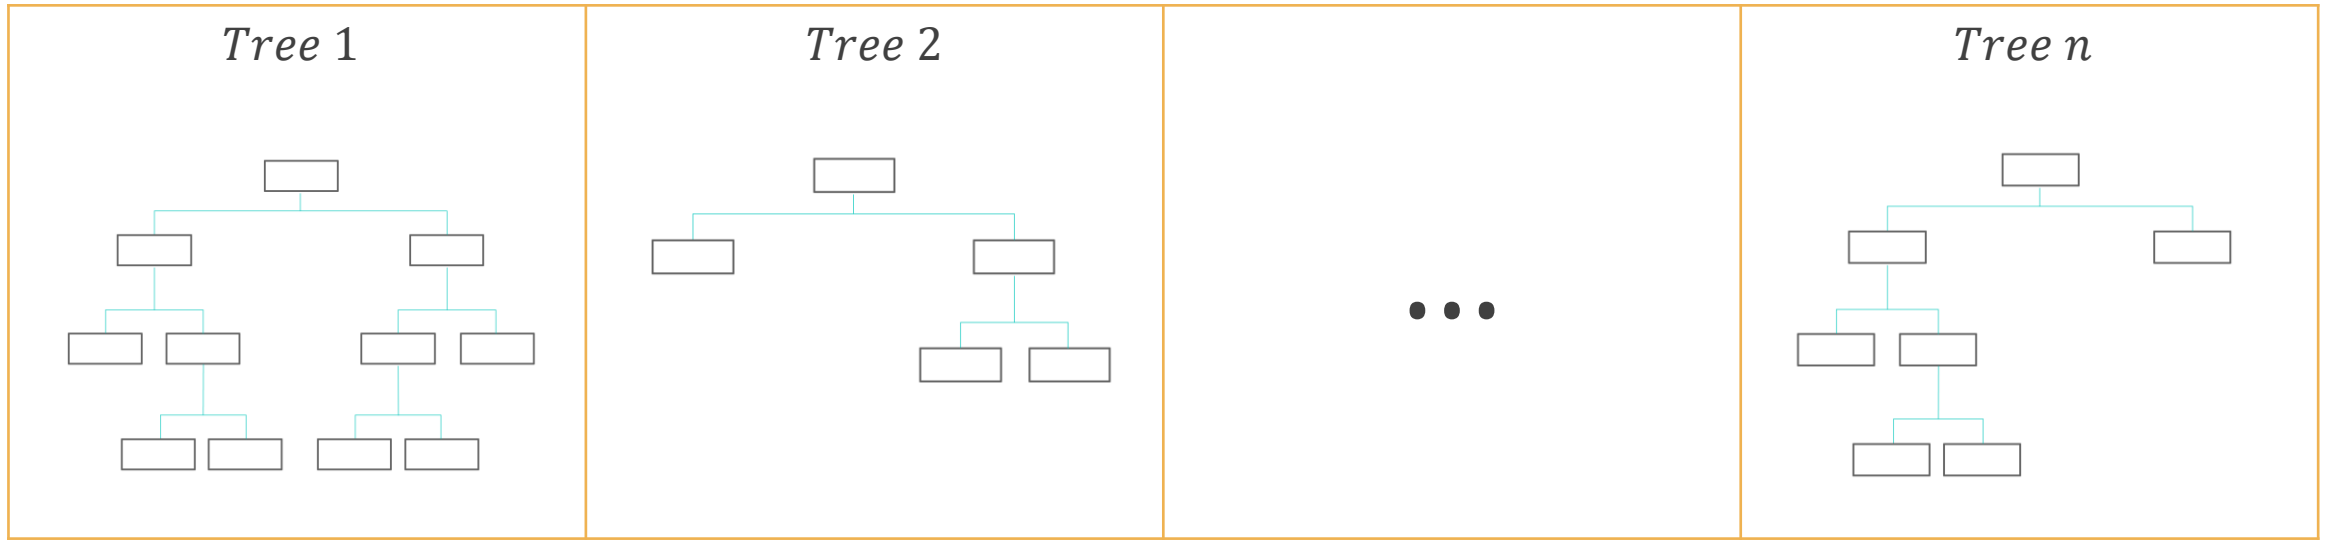
The lines connecting them show the differences between the observed and the normal.

## Wisdom of the Crowd

*Vox populi*



Sir Francis Galton



**Wisdom of the Crowd**

The final prediction (of the ensemble) is  
the averaged prediction of the individual classifiers.

A **random forest** is a **meta estimator**  
that fits a number of decision tree classifiers  
on various sub-samples of the dataset  
and uses averaging  
to **improve the predictive accuracy**  
and **control over-fitting**.

[\[scikit-learn\]](#)

# Random Forests in `scikit-learn`

```
class sklearn.ensemble.RandomForestClassifier (
```

```
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```

1

```
# Import  
from sklearn.ensemble import RandomForestClassifier
```

2

```
# Define  
clf = RandomForestClassifier()
```

3

```
# Fit  
clf.fit(x_train, y_class)
```

4

```
# Predict  
clf.predict(x_test)
```

```
class sklearn.ensemble.RandomForestClassifier (  
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```

**The function to measure  
the quality of a split.**

Supported criteria are “**gini**” for the Gini impurity and “**entropy**” for the information gain.



```
class sklearn.ensemble.RandomForestClassifier (  
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```

## The maximum depth of the tree.

If **None**, then nodes are expanded until all leaves are pure or until all leaves contain less than `min_samples_split` samples.

Recommendation: `max_depth` between 6 and 10

```
class sklearn.ensemble.RandomForestClassifier (  
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```

**The minimum number of samples  
required to split an internal node:**

If **int**, then consider **min\_samples\_split** as  
the minimum number.

If **float**, then  
**ceil(min\_samples\_split \* n\_samples)**  
are the minimum number of samples for  
each split.

Recommendation: **min\_samples\_split = 0.05**

```
class sklearn.ensemble.RandomForestClassifier (  
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```

**The minimum number of samples required to be at a leaf node.**

A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches.

Recommendation: `min_samples_leaf = 0.02`

```
class sklearn.ensemble.RandomForestClassifier (  
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```

## The number of features to consider when looking for the best split:

If **int**, then consider **max\_features** features at each split.

If **float**, then **max\_features** is a fraction and **int(max\_features \* n\_features)** features are considered at each split.

If "**auto**", then  
**max\_features=sqrt(n\_features)**.

If "**log2**", then  
**max\_features=log2(n\_features)**.

If **None**, then **max\_features=n\_features**.

Recommendation: **max\_features = 'auto'**

```
class sklearn.ensemble.RandomForestClassifier (  
    n_estimators='warn',  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features='auto',  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    bootstrap=True,  
    oob_score=False,  
    n_jobs=None,  
    random_state=None,  
    verbose=0,  
    warm_start=False,  
    class_weight=None)
```


**Whether bootstrap samples are used  
when building trees.**

Recommendation: `bootstrap = True`


# Bagging (Bootstrap Aggregation)

- **Random Forest:** For each tree, it uses a random subset of data as well as a random subset of all available features.
- The goal is to reduce the variance by doing this.

# Decision Tree



# Random Forest



# Feature Importance

```
from sklearn.datasets import make_classification

# Build a classification task using 3 informative features
X, y = make_classification(n_samples=1000,
                          n_features=10,
                          n_informative=3,
                          n_redundant=0,
                          n_repeated=0,
                          n_classes=2,
                          random_state=314,
                          shuffle=False)
```


```
# Build a forest and compute the feature importances
from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(random_state=314)

forest.fit(X, y)
```

```
importances = forest.feature_importances_

indices = np.argsort(importances)[::-1]
```





# Feature Importance


- Features used at the top of the tree contribute to the final prediction decision of a larger fraction of the input samples.
- The expected **fraction** of the samples they contribute to can thus be used as an estimate of the relative importance of the features.
- In `scikit-learn`, the **fraction of samples a feature contributes to** is combined with the **decrease in impurity** from splitting them to create a normalized estimate of the predictive power of that feature.

# Gradient Boosting Classifiers




1993





*Tree 1*




*Tree 2*



*Tree 3*





**Gradient Boosting = Boosting + Gradient Descent**

*Greedy Function Approximation: A Gradient Boosting Machine* – Jerome Friedman (1999)


*A Decision-Theoretic Generalization of On-Line Learning and an Application to Boosting* – Y. Freund, and R. Schapire (1997)



$$\boxed{F_1} = \boxed{F_0} + \boxed{\Delta_1}$$




$$F_m = F_{m-1} + \Delta_m$$



$$\hat{y} = f_0(\mathbf{x}) + \Delta_1(\mathbf{x}) + \Delta_2(\mathbf{x}) + \dots + \Delta_M(\mathbf{x})$$

MSE Loss Function



Stage $m$	Boosted Model	Model Output $\hat{y}$	Actual Error	Predicted Error $\Delta_m$
0	$F_0$	70		
1	$F_1 = F_0 + \Delta_1$	$70 + 15 = 85$	$100 - 70 = 30$	$\Delta_1 = 15$
2	$F_2 = F_1 + \Delta_2$	$85 + 20 = 105$	$100 - 85 = 15$	$\Delta_2 = 20$
3	$F_3 = F_2 + \Delta_3$	$105 - 10 = 95$	$100 - 105 = -5$	$\Delta_3 = -10$
4	$F_4 = F_3 + \Delta_4$	$95 + 5 = 100$	$100 - 95 = 5$	$\Delta_4 = 5$



**Algorithm:** *l2boost*( $X, y, M, \eta$ ) **returns** model  $F_M$

**1** Let  $F_0(X) = \frac{1}{N} \sum_{i=1}^N y_i$ , mean of target  $y$  across all observations

**2** **for**  $m = 1$  **to**  $M$  **do**

Let  $r_{m-1} = y - F_{m-1}(X)$  be the residual vector

Train **regression tree**  $\Delta_m$  on  $r_{m-1}$ , minimizing squared error

$$F_m(X) = F_{m-1}(X) + \eta \Delta_m(X)$$

**end**

**3** **return**  $F_M$

**HYPER-PARAMETERS** → **GRID SEARCH**

$M$  = number of steps  
 $\eta$  = learning rate



**Linear Regression**



**Gradient Boosting**

**Gradient Boosting `scikit-learn`**

```
class sklearn.ensemble.GradientBoostingClassifier (
```

```
    loss='deviance',
```

```
    learning_rate=0.1,
```

```
    n_estimators=100,
```

```
    subsample=1.0,
```

```
    criterion='friedman_mse',
```

```
    min_samples_split=2,
```

```
    min_samples_leaf=1,
```

```
    min_weight_fraction_leaf=0.0,
```

```
    max_depth=3,
```

```
    min_impurity_decrease=0.0,
```

```
    min_impurity_split=None,
```

```
    init=None,
```

```
    random_state=None,
```

```
    max_features=None,
```

```
    verbose=0,
```

```
    max_leaf_nodes=None,
```

```
    warm_start=False,
```

```
    presort='auto',
```

```
    validation_fraction=0.1,
```

```
    n_iter_no_change=None,
```

```
    tol=0.0001)
```

The learning rate shrinks the contribution of each tree by **learning\_rate**.

There is a trade-off between **learning\_rate** and **n\_estimators**.

$\eta$  = learning rate

```
class sklearn.ensemble.GradientBoostingClassifier (  
    loss='deviance',  
    learning_rate=0.1,  
    n_estimators=100,  
    subsample=1.0,  
    criterion='friedman_mse',  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_depth=3,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    init=None,  
    random_state=None,  
    max_features=None,  
    verbose=0,  
    max_leaf_nodes=None,  
    warm_start=False,  
    presort='auto',  
    validation_fraction=0.1,  
    n_iter_no_change=None,  
    tol=0.0001)
```

The **number of boosting stages**  
to perform.

Gradient boosting is fairly robust  
to over-fitting so a large number usually  
results in better performance.

$M$ = number of steps

```
class sklearn.ensemble.GradientBoostingClassifier (
```

```
    loss='deviance',
```

```
    learning_rate=0.1,
```

```
    n_estimators=100,
```

```
    subsample=1.0,
```

```
    criterion='friedman_mse',
```

```
    min_samples_split=2,
```

```
    min_samples_leaf=1,
```

```
    min_weight_fraction_leaf=0.0,
```

```
    max_depth=3,
```

```
    min_impurity_decrease=0.0,
```

```
    min_impurity_split=None,
```

```
    init=None,
```

```
    random_state=None,
```

```
    max_features=None,
```

```
    verbose=0,
```

```
    max_leaf_nodes=None,
```

```
    warm_start=False,
```

```
    presort='auto',
```

```
    validation_fraction=0.1,
```

```
    n_iter_no_change=None,
```

```
    tol=0.0001)
```

The **fraction of samples** to be used for fitting the individual base learners.

If smaller than 1.0 this results in  
**Stochastic Gradient Boosting.**

Choosing `subsample < 1.0` leads to  
a reduction of variance and  
an increase in bias.

```
class sklearn.ensemble.GradientBoostingClassifier (
```

```
    loss='deviance',  
    learning_rate=0.1,  
    n_estimators=100,  
    subsample=1.0,  
    criterion='friedman_mse',  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_depth=3,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    init=None,  
    random_state=None,  
    max_features=None,  
    verbose=0,  
    max_leaf_nodes=None,  
    warm_start=False,  
    presort='auto',  
    validation_fraction=0.1,  
    n_iter_no_change=None,  
    tol=0.0001)
```

The **proportion of training data** to set aside as validation set for **early stopping**.

Must be between 0 and 1.



Only used if `n_iter_no_change` is set to an integer.

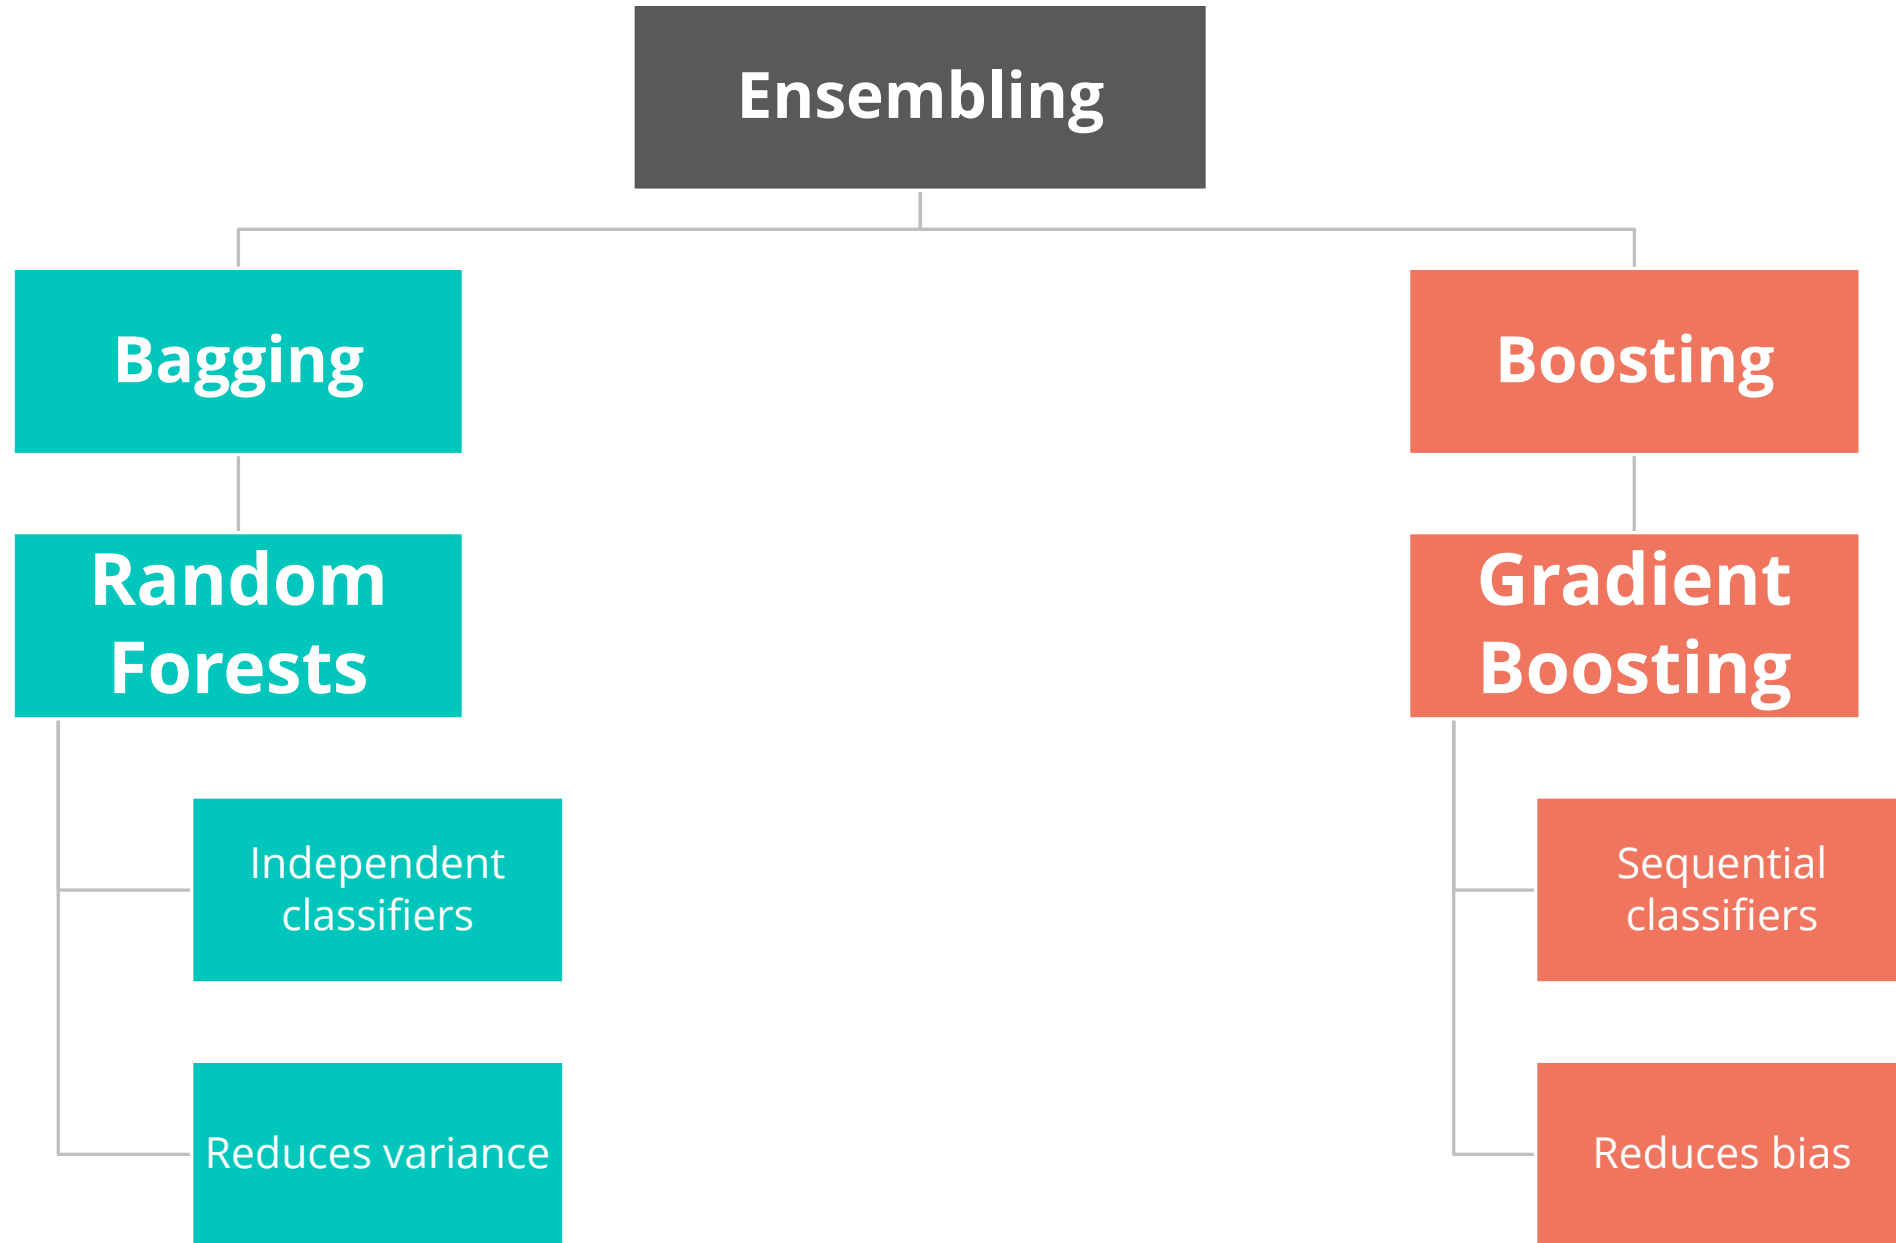
```
class sklearn.ensemble.GradientBoostingClassifier (
```

```
    loss='deviance',  
    learning_rate=0.1,  
    n_estimators=100,  
    subsample=1.0,  
    criterion='friedman_mse',  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_depth=3,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    init=None,  
    random_state=None,  
    max_features=None,  
    verbose=0,  
    max_leaf_nodes=None,  
    warm_start=False,  
    presort='auto',  
    validation_fraction=0.1,  
    n_iter_no_change=None,  
    tol=0.0001)
```

`n_iter_no_change` is used to decide  
if **early stopping** will be used  
to terminate training  
when validation score is not improving.







# Further Reading: Gradient Boosting

## Gradient boosting: Distance to target

[Terence Parr](#) and [Jeremy Howard](#)

<https://explained.ai/gradient-boosting/L2-loss.html>

### 3.2.4.3.5. `sklearn.ensemble.GradientBoostingClassifier`

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>

## Boosting algorithm: AdaBoost



SauceCat [Follow](#)

Apr 29, 2017 · 6 min read

<https://towardsdatascience.com/boosting-algorithm-adaboost-b6737a9ee60c>

## XGBoost Documentation


<https://xgboost.readthedocs.io/en/latest/>




Gradient Boosting Machine Learning  
By Trevor Hastie

# Summary


## Linear Regression



## Regression Trees





## Classification Trees



## Ensemble Trees

Random Forests & Gradient Boosting





Method	Advantages	Disadvantages
<b>Linear / Logistic Regression</b>	<ul style="list-style-type: none"> <li>○ Model fit diagnostics</li> <li>○ Interpretable coefficients, even for categorical predictors<sup>†</sup></li> <li>○ Tests for predictors</li> </ul>	<ul style="list-style-type: none"> <li>○ Simple/linear relationships</li> <li>○ Unable to handle missing values</li> <li>○ Binarize categorical predictors</li> </ul>
<b>Classification and Regression Trees</b>	<ul style="list-style-type: none"> <li>○ Categorical predictors<sup>†</sup></li> <li>○ Handles missing values and non-linear relationships</li> <li>○ Visualization/interpretation</li> <li>○ Variable importance</li> </ul>	<ul style="list-style-type: none"> <li>○ Prone to overfitting</li> <li>○ Slow for categorical data with many levels</li> </ul>
<b>Random Forest &amp; Gradient Boosting</b>	<ul style="list-style-type: none"> <li>○ Categorical predictors<sup>†</sup> &amp; missing values</li> <li>○ Variable importance</li> <li>○ Controls overfitting (bias and/or variance)</li> </ul>	<ul style="list-style-type: none"> <li>○ No visualization</li> <li>○ Slow for large datasets</li> <li>○ Hyper-parameters is required</li> </ul>
<b>Support Vector Machine</b>	<ul style="list-style-type: none"> <li>○ Vary complexity by changing kernel/tuning parameters</li> </ul>	<ul style="list-style-type: none"> <li>○ Hard to visualize/interpret</li> <li>○ Hyper-parameters is required</li> <li>○ Binarize categorical predictors</li> </ul>
<b>Neural Network</b>	<ul style="list-style-type: none"> <li>○ Vary complexity by changing number of layers/tuning parameters</li> <li>○ High accuracy</li> </ul>	<ul style="list-style-type: none"> <li>○ Hard to visualize/interpret, “black box”</li> <li>○ Tuning parameters, stopping criteria</li> <li>○ Binarize categorical predictors</li> </ul>